

Advanced features in INLA and inlabru

University of Zurich, March, 2022

Instructor: Sara Martino

Department of Mathematical Science (NTNU)



NTNU

Norwegian University of
Science and Technology

What have we learned

Spatial models

Gaussian Random Field models

The SPDE approach

Fitting spatial models

Several likelihood

What have we learned	Spatial models	Gaussian Random Field models	The SPDE approach	Fitting spatial
oooooooo	ooooo	oooo	ooo	oooooooooooooooo

What have we learned

INLA in a nutshell

- many data sets these days are complex, resulting in complex models, e.g. with complex dependence structures (spatial, temporal, etc..)
- usually Markov chain Monte Carlo (MCMC) methods have been used to fit these models
 - (realistically) complex models result in very long running times
 - often impossible (or unrealistic) to fit
- INLA (Integrated nested Laplace approximation) is an alternative to MCMC
 - much, much **faster**
 - suitable for a specific (but very large!) class of models

INLA in a nutshell

Three main ingredients in INLA

- Gaussian Markov random fields
- Latent Gaussian models
- Laplace approximations

which together (with a few other things) give a very nice tool for Bayesian inference

- quick
- accurate

Models amenable to INLA: Latent Gaussian Models

$$\mathbf{y}|\mathbf{x}, \theta \sim \prod_i \pi(y_i|x_i, \theta)$$

Likelihood

$$\mathbf{x}|\theta \sim \exp\left(-\frac{1}{2}\mathbf{x}^T\mathbf{Q}(\theta)\mathbf{x}\right)$$

Latent field (GMRF)

$$\theta \sim \pi(\theta)$$

Hyperparameter

- Each data point depends on only one of the elements in the latent Gaussian field \mathbf{x} .

Models amenable to INLA: Latent Gaussian Models

$$\mathbf{y}|\mathbf{x}, \theta \sim \prod_i \pi(y_i|x_i, \theta) \quad \text{Likelihood}$$

$$\mathbf{x}|\theta \sim \exp\left(-\frac{1}{2}\mathbf{x}^T\mathbf{Q}(\theta)\mathbf{x}\right) \quad \text{Latent field (GMRF)}$$

$$\theta \sim \pi(\theta) \quad \text{Hyperparameter}$$

- Each data point depends on only one of the elements in the latent Gaussian field \mathbf{x} .
- The size of the hyperparameter vector θ is small (say < 15)

Models amenable to INLA: Latent Gaussian Models

$$\mathbf{y}|\mathbf{x}, \theta \sim \prod_i \pi(y_i|x_i, \theta)$$

Likelihood

$$\mathbf{x}|\theta \sim \exp\left(-\frac{1}{2}\mathbf{x}^T\mathbf{Q}(\theta)\mathbf{x}\right)$$

Latent field (GMRF)

$$\theta \sim \pi(\theta)$$

Hyperparameter

- Each data point depends on only one of the elements in the latent Gaussian field \mathbf{x} .
- The size of the hyperparameter vector θ is small (say < 15)
- The latent field \mathbf{x} , can be large but it is endowed with some conditional independence (Markov) properties so that the precision matrix $Q(\theta)$ is sparse.

Models amenable to INLA: Latent Gaussian Models

$$\mathbf{y}|\mathbf{x}, \theta \sim \prod_i \pi(y_i|x_i, \theta)$$

Likelihood

$$\mathbf{x}|\theta \sim \exp\left(-\frac{1}{2}\mathbf{x}^T\mathbf{Q}(\theta)\mathbf{x}\right)$$

Latent field (GMRF)

$$\theta \sim \pi(\theta)$$

Hyperparameter

- Each data point depends on only one of the elements in the latent Gaussian field \mathbf{x} .
- The size of the hyperparameter vector θ is small (say < 15)
- The latent field \mathbf{x} , can be large but it is endowed with some conditional independence (Markov) properties so that the precision matrix $Q(\theta)$ is sparse.
- The linear predictor depends **linearly** on the unknown smooth function of covariates.

Models amenable to INLA: Latent Gaussian Models

$$\mathbf{y}|\mathbf{x}, \theta \sim \prod_i \pi(y_i|x_i, \theta)$$

Likelihood

$$\mathbf{x}|\theta \sim \exp\left(-\frac{1}{2}\mathbf{x}^T\mathbf{Q}(\theta)\mathbf{x}\right)$$

Latent field (GMRF)

$$\theta \sim \pi(\theta)$$

Hyperparameter

- Each data point depends on only one of the elements in the latent Gaussian field \mathbf{x} .
- The size of the hyperparameter vector θ is small (say < 15)
- The latent field \mathbf{x} , can be large but it is endowed with some conditional independence (Markov) properties so that the precision matrix $Q(\theta)$ is sparse.
- The linear predictor depends **linearly** on the unknown smooth function of covariates.

All of this is implemented in the INLA library in R.

Yesterday we saw how to implement simple models with INLA

There are many more features that can be implemented with
INLA

All of this is implemented in the INLA library in R.

Yesterday we saw how to implement simple models with INLA

There are many more features that can be implemented with INLA

- Spatial models
- Use several likelihoods
- Predict linear combinations of elements of the latent field
- Group/copy/replicate random fields

All of this is implemented in the INLA library in R.

Yesterday we saw how to implement simple models with INLA

There are many more features that can be implemented with INLA

- Spatial models
- Use several likelihoods
- Predict linear combinations of elements of the latent field
- Group/copy/replicate random fields

Many of these features are **a lot** easier to implement in `inlabru`

inlabru in a nutshell

inlabru is a friendlier version of R-INLA

- it makes INLA more accessible to the user
- makes complex features and predictions (especially for spatial data) a lot easier
- is a softer-wrapper around INLA
- allows to release the

inlabru

- Installation:
 - There is a CRAN version

```
install.packages("inlabru")
```

- You can also install the development version of inlabru from GitHub (recommended)

```
# install.packages("remotes")  
remotes::install_github("inlabru-org/inlabru",  
  ref="devel")
```

- Documentation
 - Web site: <https://sites.google.com/inlabru.org/inlabru>
 - Github: <https://github.com/inlabru-org/inlabru>

What have we learned

○○○○○○○

Spatial models

●○○○○

Gaussian Random Field models

○○○○

The SPDE approach

○○○

Fitting spatial

○○○○○○○○○○○○○

Spatial models

Types of Spatial Data


We can distinguish three types of spatial data

- Discrete space
 - data on a spatial grid
- Continuous space:
 - geostatistical data
 - spatial point data

Discrete Counts:

Data on a spatial grid

- examples: number of individuals in a region, average rainfall in a province
- (originally geostatistical or point data; gridded for practical reasons)




Observed response(s):

- Measurement over each grid cell (e.g. number of individuals in cell; rainfall in province)

Continuous Space: Geostatistics

- phenomenon that is continuous in space
- examples: nutrient levels in soil, salinity in the sea
- measurements at a given set of locations that are determined by surveyor




Observed response(s):

- measurement(s) taken at given locations

Continuous Space: Point Process

- locations of objects (individuals) in space (typically 2D)
- examples: locations of trees in a forest, groups of animals



Observed response(s):

- x, y coordinates of points (individuals/groups)
- maybe also properties of individuals/groups (“marks”)

Gaussian Random Field models

Gaussian Random fields

Definition: A random function $u(x) : R^d \rightarrow R$ is a Gaussian random field if for any finite collection of locations, (x_1, \dots, x_n) , $x_i \in R^d$, the joint distribution of $\mathbf{u} = (u(x_1), \dots, u(x_n))$ is $\mathbf{u} \sim N(0, \Sigma)$, and

$$E(u(x)) = 0,$$

$$Cov(u(x), u(x')) = R(x, x'), \quad \Sigma_{ij} = R(x_i, x_j)$$

for some expectation function $\mu(\cdot)$ and positive definite covariance function $R(\cdot, \cdot)$. Σ is the covariance matrix for the specific location collection.

What have we learned
○○○○○○○


Spatial models
○○○○○

Gaussian Random Field models
○○●○

The SPDE approach
○○○

Fitting spatial
○○○○○○○○○○○○○○

Example



Gaussian Random field

GRF are a very popular model

- Flexible and easy to use
- Can be part of the latent Gaussian field in a LGM


However:

- **computationally inefficient** (the precision matrix is dense)
- not flexible enough (complicated boundary, barrier, . . .)

The SPDE approach


The SPDE approach

- Matern fields can be seen as solution to a PDE
- Using finite element methods such solution can be represented using a GRMF



Advantages of the SPDE approach

- Computationally fast
- Allows for flexible modeling
 - non-stationary models (anisotropy)
 - models on a sphere
 - non separable models




All these models (and my more) can be fitted with R-INLA and `inlabru`

Fitting spatial models

Example: Meuse Data

Measures of zinc concentration.



The model

$$\begin{aligned}\log(Y(s)) &\sim \mathcal{N}(\eta(s), \sigma_y^2) \\ \eta(s) &= \alpha + u(s)\end{aligned}$$


where

- $Y(s)$ is the measure of zinc in location s
- α a common intercept
- $u(s)$ the Matern Gaussian field

Define the SPDE representation: The mesh

1. Define the mesh

```
mesh <- inla.mesh.2d(loc.domain = cbind(meuse$x, meuse$y),  
                     max.edge = c(150, 500),  
                     offset = c(100, 250) )
```



Define the SPDE representation: The mesh

- All random field models need to be discretised for practical calculations.
- The SPDE models were developed to provide a consistent model definition across a range of discretisations.
- We use finite element methods with local, piecewise linear basis functions defined on a triangulation of a region of space containing the domain of interest.
- Deviation from stationarity is generated near the boundary of the region.
- The choice of region and choice of triangulation affects the numerical accuracy.

Define the SPDE representation: The mesh

- All random field models need to be discretised for practical calculations.
- The SPDE models were developed to provide a consistent model definition across a range of discretisations.
- We use finite element methods with local, piecewise linear basis functions defined on a triangulation of a region of space containing the domain of interest.
- Deviation from stationarity is generated near the boundary of the region.
- The choice of region and choice of triangulation affects the numerical accuracy.

Two separate issues:

- Continuous space with bounded domain: Boundary effect
- Discretised model: Numerical accuracy

Sometimes the boundary effect may be desirable.

Define the SPDE representation: The mesh

- Too fine meshes \rightarrow heavy computation
- Too coarse mesh \rightarrow not accurate enough

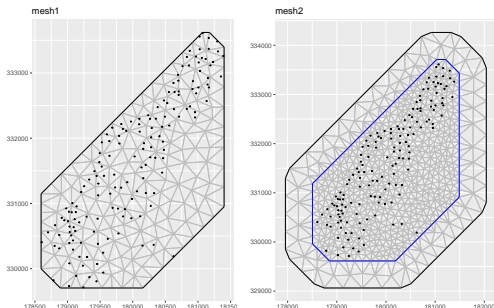
Some guidelines

- Create triangulation meshes with `inla.mesh.2d()`
- Move undesired boundary effects away from the domain of interest by extending to a smooth external boundary (`inla.nonconvex.hull(loc, convex)`, `convex \geq correlation range`)
- Use a coarser resolution in the extension to reduce computational cost (`max.edge=c(inner, outer)`)
- Use a fine resolution (subject to available computational resources) for the domain of interest (inner correlation range) and filter out small input point clusters (`0 < cutoff < inner`)
- Coastlines and similar can be added to the domain specification in `inla.mesh.2d()`

Define the SPDE representation: The mesh

```
mesh1 = inla.mesh.2d(loc.domain = cbind(meuse$x, meuse$y),
                     max.edge = 350,
                     offset = 10)
```

```
mesh2 = inla.mesh.2d(loc.domain = cbind(meuse$x, meuse$y),
                     max.edge = c(150, 500),
                     cutoff = 100,
                     offset = c(100, 550) )
```



Define the SPDE representation: The SPDE model

```
meuse.spde <- inla.spde2.pcmatern(mesh = mesh,  
                                   prior.sigma = c(1, 0.1),  
                                   prior.range = c(1000, 0.5))
```

PC-priors for the range ρ and the standard deviation σ

- Define the prior for the range `prior.range` =
(range0, Prange) $\text{Prob}(\rho < \rho_0) = p_\rho$
- Define the prior for the range `prior.sigma` =
(sigma0, Psigma) $\text{Prob}(\sigma < \sigma_0) = p_\sigma$

Run the model inlabru

```

# create a spatial object
coordinates(meuse) = c("x","y")
# covariate values
dist_SPDE = SpatialPixelsDataFrame(data$dist_raster[,c(1,2)],
                                   data = data.frame(dist = data$dist_raster[,3]))

# model components
cmp = ~ Intercept(1) + dist(dist_SPDE, model = "linear") +
      spde(coordinates, model = meuse.spde)
# define likelihood
lik = like(formula = Y ~ Intercept + dist + spde,
           family = "gaussian",
           data = meuse)

#fit the model
fit <- bru(cmp, lik)
# define prediction area
pix <- pixels(mesh, nx = 200, ny = 200, mask = boundary)
# generate predictions
pred = predict(fit, pix, ~ data.frame(
                                   spde = spde,
                                   logscale = Intercept + dist + spde,
                                   naturalscale = exp(Intercept + dist + spde)))

```

Notes!

- The data are a spatial object!
- The covariates are stored in a `SpatialPixelsDataFrame` and need to cover all the mesh nodes

What have we learned
○○○○○○○



Spatial models
○○○○○

Gaussian Random Field models
○○○○

The SPDE approach
○○○

Fitting spatial
○○○○○○○○○○○●

Predictions: The SPDE field



What have we learned

oooooooo

Spatial models

ooooo

Gaussian Random Field models

oooo



The SPDE approach

ooo

Fitting spatial

oooooooooooooooo

Predictions: The log concentrations



What have we learned

oooooooo

Spatial models

ooooo

Gaussian Random Field models

oooo



The SPDE approach

ooo

Fitting spatial

oooooooooooooooo

Predictions: The concentrations



Same in plain INLA (1)

```

A.meuse <- inla.spde.make.A(mesh = mesh, loc = coordinates(meuse))
s.index <- inla.spde.make.index(name = "spatial.field",
  n.spde = meuse.spde$n.spde)

#Create data structure
meuse.stack <- inla.stack(data = list(zinc = meuse$zinc),
  A = list(A.meuse, 1),
  effects = list(c(s.index, list(Intercept = 1)),
    list(dist = meuse$dist)),
  tag = "meuse.data")

data(meuse.grid)
coordinates(meuse.grid) = ~x+y
gridded(meuse.grid) = TRUE

#Create data structure for prediction
A.pred <- inla.spde.make.A(mesh = mesh, loc = coordinates(meuse.grid))
meuse.stack.pred <- inla.stack(data = list(zinc = NA),
  A = list(A.pred, 1),
  effects = list(c(s.index, list(Intercept = 1)),
    list(dist = meuse.grid$dist)),
  tag = "meuse.pred")

#Join stack
join.stack <- inla.stack(meuse.stack, meuse.stack.pred)

```

Same in plain INLA (2)

```
#Fit model
form <- log(zinc) ~ -1 + Intercept + dist + f(spatial.field, model = spde)

m1 <- inla(form, data = inla.stack.data(join.stack, spde = meuse.spde),
  family = "gaussian",
  control.predictor = list(A = inla.stack.A(join.stack), compute = TRUE))
```

Note: We still have not compute predictions...and this is not too easy in plain INLA!

When is `inlabru` easier to use

- spatial modeling.
- point processes.
- multiple likelihoods
- when interested in spatial predictions

When is `inlabru` easier to use

- spatial modeling.
- point processes.
- multiple likelihoods
- when interested in spatial predictions
- `inlabru` is also useful if one has non-linearities in the predictor η
 - born for ecological models (for example transect sampling)
but used also in other fields

What have we learned

oooooooo

Spatial models

ooooo

Gaussian Random Field models

oooo

The SPDE approach

ooo

Fitting spatial

oooooooooooooooo

Several likelihood

Example: Coregionalization model

$$y_1(s) = \alpha_1 + u(s) + e_1(s)$$

$$y_2(s) = \alpha_2 + \lambda u(s) + e_2(s)$$

where the α_k are intercepts, $u(s)$ is the spatial effect, λ is a weights for spatial effects and $e_k(s)$ are uncorrelated error terms, with $k = 1, 2, 3$.