

Getting INLA	Implementing the INLA algorithm	How to use INLA	Simple example	Add random effects	Prediction
0000	000000	0000000	000000000000	000000	0000000

Bayesian Statistics with R-INLA

University of Zurich, March, 2022

Instructor: Sara Martino

Department of Mathematical Science (NTNU)



Notes

Getting INLA	Implementing the INLA algorithm	How to use INLA	Simple example	Add random effects	Prediction	Smoothing binary time series	Disease Mapping	Changing the prior	Repeated Poisson counts	Control statements
0000	000000	0000000	000000000000	000000	0000000	0000000	00000000	000	0000000000000000	000

Getting INLA

Implementing the INLA algorithm

How to use INLA

Simple example

Add random effects

Prediction

Smoothing binary time series

Disease Mapping

Changing the prior

Repeated Poisson counts

Control statements

Notes

Getting INLA	Implementing the INLA algorithm	How to use INLA	Simple example	Add random effects	Prediction	Smoothing binary time series	Disease Mapping	Changing the prior	Repeated Poisson counts	Control statements
0000	000000	0000000	000000000000	000000	0000000	0000000	00000000	000	0000000000000000	000

Notes

Getting INLA oooo	Implementing the INLA algorithm oooooo	How to use INLA ooooooo	Simple example oooooooooooo	Add random effects oooooo	Prediction ooooooo	Smoothing binary time series ooooooo	Disease Mapping ooooooooo	Changing the prior ooo	Repeated Poisson counts oooooooooooooooooooo	Control statements ooo
----------------------	---	----------------------------	--------------------------------	------------------------------	-----------------------	---	------------------------------	---------------------------	---	---------------------------

What have we learned in the morning...

- What is a LGM
- Which kind of models are amenable to INLA
- How does INLA work...

Notes

Getting INLA oooo	Implementing the INLA algorithm oooooo	How to use INLA ooooooo	Simple example oooooooooooo	Add random effects oooooo	Prediction ooooooo	Smoothing binary time series ooooooo	Disease Mapping ooooooooo	Changing the prior ooo	Repeated Poisson counts oooooooooooooooooooo	Control statements ooo
----------------------	---	----------------------------	--------------------------------	------------------------------	-----------------------	---	------------------------------	---------------------------	---	---------------------------

What have we learned in the morning...

- What is a LGM
- Which kind of models are amenable to INLA
- How does INLA work...
- ..you have even implemented it yourself! :-)

Notes

Getting INLA oooo	Implementing the INLA algorithm oooooo	How to use INLA ooooooo	Simple example oooooooooooo	Add random effects oooooo	Prediction ooooooo	Smoothing binary time series ooooooo	Disease Mapping ooooooooo	Changing the prior ooo	Repeated Poisson counts oooooooooooooooooooo	Control statements ooo
----------------------	---	----------------------------	--------------------------------	------------------------------	-----------------------	---	------------------------------	---------------------------	---	---------------------------

Good News!

All the theory we have seen is wrapped up in the R-package **INLA** which is easy to use.

Notes

Getting INLA ●○○○	Implementing the INLA algorithm ○○○○○	How to use INLA ○○○○○○	Simple example ○○○○○○○○○○	Add random effects ○○○○○	Prediction ○○○○○○	Smoothing binary time series ○○○○○○	Disease Mapping ○○○○○○○	Changing the prior ○○○	Repeated Poisson counts ○○○○○○○○○○○○○○	Control statements ○○○
----------------------	--	---------------------------	------------------------------	-----------------------------	----------------------	--	----------------------------	---------------------------	---	---------------------------

Getting INLA

Notes

Getting INLA ○○●○	Implementing the INLA algorithm ○○○○○	How to use INLA ○○○○○○	Simple example ○○○○○○○○○○	Add random effects ○○○○○	Prediction ○○○○○○	Smoothing binary time series ○○○○○○	Disease Mapping ○○○○○○○	Changing the prior ○○○	Repeated Poisson counts ○○○○○○○○○○○○○○	Control statements ○○○
----------------------	--	---------------------------	------------------------------	-----------------------------	----------------------	--	----------------------------	---------------------------	---	---------------------------

Getting INLA

Notes

- The web page www.r-inla.org contains source-code, worked-through examples, reports and instructions for installing the package.

Getting INLA ○○●○	Implementing the INLA algorithm ○○○○○	How to use INLA ○○○○○○	Simple example ○○○○○○○○○○	Add random effects ○○○○○	Prediction ○○○○○○	Smoothing binary time series ○○○○○○	Disease Mapping ○○○○○○○	Changing the prior ○○○	Repeated Poisson counts ○○○○○○○○○○○○○○	Control statements ○○○
----------------------	--	---------------------------	------------------------------	-----------------------------	----------------------	--	----------------------------	---------------------------	---	---------------------------

Getting INLA

- The R-package INLA works on Linux, Windows and Mac and can be installed within R by

```
# stable version
install.packages("INLA",
  repos=c(getOption("repos"),
    INLA="https://inla.r-inla-download.org/R/stable"),
  dep=TRUE)

# devel version
install.packages("INLA",
  repos=c(getOption("repos"),
    INLA="https://inla.r-inla-download.org/R/testing"),
  dep=TRUE)
```

and then upgraded in R as:

```
inla.upgrade(testing = TRUE)
```

****NB** You need R version 4.1 or newer!!**

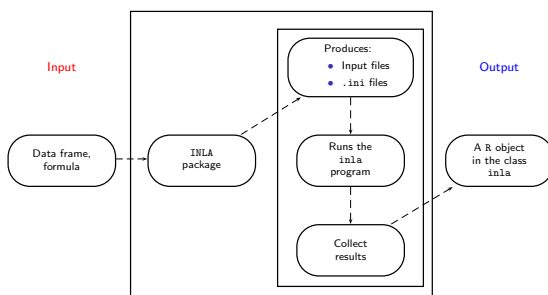
Notes

```
## R-INLA version .....: 22.02.16-2 _____  
## Date .....: Wed Feb 16 02:22:38 PM +03 2022 (Version_22.02.16-2)  
## Maintainers .....: Havard Rue <hrue@r-inla.org>  
## : Finn Lindgren <finn.lindgren@gmail.com>  
## : Elias Teixeira Krainski <elias@r-inla.org>  
## Main web-page .....: www.r-inla.org _____  
## Download-page .....: inla.r-inla-download.org  
## Repository .....: github.com/hrue/r-inla _____  
## Email support .....: help@r-inla.org  
## : r-inla-discussion-group@googlegroups.com
```

```

version_22.02.16-2)
com>
la.org>
s.com

```

[illegible]

Getting INLA
oooo

Implementing the INLA algorithm
ooo●ooo

How to use INLA
oooooooo

Simple example
oooooooooooo

Add random effects
oooooo

Prediction
oooooooo

Smoothing binary time series
oooooooo

Disease Mapping
oooooooo

Changing the prior
ooo

Repeated Poisson counts
oooooooooooooooo

Control statements
ooo

What happens in the black box?

The implementation of the INLA method consists of three parts:

- **GMRFLib-Library:** A library for GMRFs written in **C**
- **inla-program:** The implementation of INLA written in **C**
- **INLA package for R:** An R-interface to the **inla**-program

The first two are *not* particularly user-friendly. They are used in the background by the **INLA** package.

Notes

Getting INLA
oooo

Implementing the INLA algorithm
ooo●ooo

How to use INLA
oooooooo

Simple example
oooooooooooo

Add random effects
oooooo

Prediction
oooooooo

Smoothing binary time series
oooooooo

Disease Mapping
oooooooo

Changing the prior
ooo

Repeated Poisson counts
oooooooooooooooo

Control statements
ooo

Implementing INLA

All procedures required to perform INLA need to be carefully implemented to achieve a good speed; easier to implement a slow version of INLA.

Notes

Getting INLA
oooo

Implementing the INLA algorithm
ooo●ooo

How to use INLA
oooooooo

Simple example
oooooooooooo

Add random effects
oooooo

Prediction
oooooooo

Smoothing binary time series
oooooooo

Disease Mapping
oooooooo

Changing the prior
ooo

Repeated Poisson counts
oooooooooooooooo

Control statements
ooo

Implementing INLA

All procedures required to perform INLA need to be carefully implemented to achieve a good speed; easier to implement a slow version of INLA.

- **The GMRFLib-library**
 - Basic library written in C, user friendly for programmers

Notes

Getting INLA
○○○○

Implementing the INLA algorithm
○○○○●

How to use INLA
○○○○○○○

Simple example
○○○○○○○○○○○○○

Add random effects
○○○○○○○

Prediction
○○○○○○○

Smoothing binary time series
○○○○○○○

Disease Mapping
○○○○○○○○○

Changing the prior
○○○

Repeated Poisson counts
○○○○○○○○○○○○○○○○○

Control statements
○○○

Implementing INLA

All procedures required to perform INLA need to be carefully implemented to achieve a good speed; easier to implement a slow version of INLA.

- The **GMRFLib**-library
- The **inla**-program
 - Define *latent Gaussian models* and interface with the GMRFLib-library
 - Avoids the need for C-programming
 - Models are defined using **.ini**-files
 - Requires to write input files in a special format
 - **inla**-program write all the results (E/Var/marginals) to files

Notes

Getting INLA
○○○○

Implementing the INLA algorithm
○○○○●

How to use INLA
○○○○○○○

Simple example
○○○○○○○○○○○○○

Add random effects
○○○○○○○

Prediction
○○○○○○○

Smoothing binary time series
○○○○○○○

Disease Mapping
○○○○○○○○○

Changing the prior
○○○

Repeated Poisson counts
○○○○○○○○○○○○○○○○○

Control statements
○○○

Implementing INLA

All procedures required to perform INLA need to be carefully implemented to achieve a good speed; easier to implement a slow version of INLA.

- The **GMRFLib**-library
- The **inla**-program
- The **INLA package for R**
 - R-interface to the **inla**-program. (That's why its not on CRAN.)
 - Convert **formula**-statements into **.ini**-files definitions
 - It also does much more (for example for survival models or when using **inlabru**)

Notes

Getting INLA
○○○○

Implementing the INLA algorithm
○○○○○

How to use INLA
●○○○○○○

Simple example
○○○○○○○○○○○○○

Add random effects
○○○○○○○

Prediction
○○○○○○○

Smoothing binary time series
○○○○○○○

Disease Mapping
○○○○○○○○○

Changing the prior
○○○

Repeated Poisson counts
○○○○○○○○○○○○○○○○○

Control statements
○○○

How to use INLA

Notes

Getting INLA
0000

Implementing the INLA algorithm
000000

How to use INLA
0●000000

Simple example
000000000000

Add random effects
000000

Prediction
0000000

Smoothing binary time series
0000000

Disease Mapping
00000000

Changing the prior
000

Repeated Poisson counts
0000000000000000

Control statements
000

How to use INLA

There are essentially four parts to an INLA-program:

1. **Data organisation:** Make an object to store response, covariates,

```
data = data.frame(y = y, x = x)
```
2. **Use the ‘formula’-notation** to specify the model (similar to `lm` and `glm` functions)

```
formula = y~x
```
3. **Call the ‘inla’-program**

```
res = inla(formula, data=data, family="gaussian")
```
4. **Extract posterior information**, e.g.for a first overview use

```
summary(res)
```

Notes

Getting INLA
0000

Implementing the INLA algorithm
000000

How to use INLA
0●000000

Simple example
000000000000

Add random effects
000000

Prediction
0000000

Smoothing binary time series
0000000

Disease Mapping
00000000

Changing the prior
000

Repeated Poisson counts
0000000000000000

Control statements
000

Data organization

The responses and covariates are collected in a **list or data frame**. Assume response `y`, covariates `x1` and `x2`, and time index `t`. Then they can be organized with:

```
# Option 1  
  
data = list(y = y, x1 = x1, x2 = x2, t = t)  
  
# Option 2  
  
data = data.frame(y = y, x1 = x1, x2 = x2, t = t)
```

Notes

Getting INLA
0000

Implementing the INLA algorithm
000000

How to use INLA
0●000000

Simple example
000000000000

Add random effects
000000

Prediction
0000000

Smoothing binary time series
0000000

Disease Mapping
00000000

Changing the prior
000

Repeated Poisson counts
0000000000000000

Control statements
000

formula: specifying the linear predictor

The model is specified through a ‘formula’ similar to `glm`:

```
formula = y ~ x1 + x2 + f(t, ...)
```

- `y` is the name of the response in the `data` object
- The fixed effects are given i.i.d. Gaussian priors
- The `f()` function specifies random effects (e.g. temporal, spatial, smooth effect of covariates and Besag model)
- Use `-1` in the formula if you don’t want an automatic intercept

Notes

Getting INLA
○○○○

Implementing the INLA algorithm
○○○○○○

How to use INLA
○○○○●○○

Simple example
○○○○○○○○○○○○

Add random effects
○○○○○○

Prediction
○○○○○○○

Smoothing binary time series
○○○○○○○

Disease Mapping
○○○○○○○○

Changing the prior
○○○

Repeated Poisson counts
○○○○○○○○○○○○○○○○

Control statements
○○○

The inla() function

```
result = inla(  
  # Description of linear predictor  
  formula,  
  # Likelihood  
  family = "gaussian",  
  # List or data frame with response,  
  # covariates, etc.  
  data = data,  
  ## This is all that is needed for a basic call  
  
  ## # check what happens  
  verbose = TRUE,  
  # ,..., there are also some "control statements"  
  # to customize things  
  # This you need if you later want to sample from the  
  # fitted model  
  control.compute=list(config = TRUE)  
)
```

Notes

Getting INLA
○○○○

Implementing the INLA algorithm
○○○○○○

How to use INLA
○○○○○●○

Simple example
○○○○○○○○○○○○

Add random effects
○○○○○○

Prediction
○○○○○○○

Smoothing binary time series
○○○○○○○

Disease Mapping
○○○○○○○○

Changing the prior
○○○

Repeated Poisson counts
○○○○○○○○○○○○○○○○

Control statements
○○○

Likelihood functions

- gaussian
- T
- poisson
- nbinomial
- binomial
- exponential
- weibull
- gev
- coxph

For a complete list type

```
names(inla.models())$likelihood)
```

Notes

Getting INLA
○○○○

Implementing the INLA algorithm
○○○○○○

How to use INLA
○○○○○○●

Simple example
○○○○○○○○○○○○

Add random effects
○○○○○○

Prediction
○○○○○○○

Smoothing binary time series
○○○○○○○

Disease Mapping
○○○○○○○○

Changing the prior
○○○

Repeated Poisson counts
○○○○○○○○○○○○○○○○

Control statements
○○○

Posterior inference

Main functions:

```
# look at a first summary  
summary(result)  
# plot the main results  
# (does not use ggplot...)  
plot(result)  
# rerun the model to get better  
# estimate of the hyperparameters  
result2 = inla.hyperpar(result)  
# sample from the fitted model  
# this can be very useful sometimes!  
sample = inla.posterior.sample(results)
```

Notes

Getting INLA	Implementing the INLA algorithm	How to use INLA	Simple example	Add random effects	Prediction
0000	000000	0000000	●0000000000	000000	0000000

Simple example

Notes

Getting INLA	Implementing the INLA algorithm	How to use INLA	Simple example	Add random effects	Prediction
0000	000000	0000000	●0000000000	000000	0000000

Example: Simple linear regression

- **Stage 1:** Gaussian likelihood

$$y_i|\eta_i \sim \mathcal{N}(\eta_i, \sigma^2)$$

- **Stage 2:** Covariates are connected to likelihood by

$$\eta_i = \beta_0 + \beta_1 x_i$$

- **Stage 3:** σ^2 : variance of observation noise

Notes

Getting INLA	Implementing the INLA algorithm	How to use INLA	Simple example	Add random effects	Prediction
0000	000000	0000000	●0000000000	000000	0000000

Example: Simple linear regression

```
# Generate data
x = runif(10)
y = 1 + 2*x + rnorm(n = 100, sd = 0.1)

# Run inla
formula = y ~ 1 + x
result = inla(formula,
               data = data.frame(x = x, y = y),
               family = "gaussian")
```

Notes

Organization of the inla-object

```
names(result)

## [1] "names.fixed"          "summary.fixed"
## [3] "marginals.fixed"      "summary.lincomb"
## [5] "marginals.lincomb"    "size.lincomb"
## [7] "summary.lincomb.derived" "marginals.lincomb.derived"
## [9] "size.lincomb.derived" "mlik"
## [11] "cpo"                  "po"
## [13] "vaic"                 "model.random"
## [15] "summary.random"       "marginals.random"
## [17] "size.random"          "summary.linear.predictor"
## [19] "marginals.linear.predictor" "summary.fitted.values"
## [21] "marginals.fitted.values" "size.linear.predictor"
## [23] "summary.hyperpar"      "marginals.hyperpar"
## [25] "internal.summary.hyperpar" "internal.marginals.hyperpar"
## [27] "offset.linear.predictor" "model.spde2.blc"
## [29] "summary.spde2.blc"     "marginals.spde2.blc"
## [31] "size.spde2.blc"       "model.spde3.blc"
## [33] "summary.spde3.blc"    "marginals.spde3.blc"
## [35] "size.spde3.blc"      "logfile"
## [37] "misc"                 "dic"
## [39] "mode"                 "joint.hyper"
## [41] "nhyper"               "version"
## [43] "qi"                   "graph"
## [45] "ok"                   "cpu.used"
## [47] "all.hyper"            ".args"
## [49] "call"                 "model.matrix"
```

Notes

Organization of the inla-object

You can find summary information in

```
## [1] "summary.fixed"          "summary.lincomb"
## [3] "summary.lincomb.derived" "summary.random"
## [5] "summary.linear.predictor" "summary.fitted.values"
## [7] "summary.hyperpar"      "internal.summary.hyperpar"
## [9] "summary.spde2.blc"     "summary.spde3.blc"
```

for example

```
result$summary.fixed

##              mean          sd 0.025quant  0.5quant 0.975quant      mode
## (Intercept) 0.9886826 0.02279676   0.943834 0.9886819   1.033493 0.9886825
## x           2.0286140 0.04573803   1.938632 2.0286128   2.118518 2.0286143
##              kld
## (Intercept) 3.088291e-06
## x           3.088224e-06
```

Notes

Organization of the inla-object

You can find estimated posterior marginals in

```
## [1] "marginals.fixed"          "marginals.lincomb"
## [3] "marginals.lincomb.derived" "marginals.random"
## [5] "marginals.linear.predictor" "marginals.fitted.values"
## [7] "marginals.hyperpar"      "internal.marginals.hyperpar"
## [9] "marginals.spde2.blc"     "marginals.spde3.blc"
```

Each object is thereby a list. Get the marginal for intercept:

```
head(result$marginals.fixed[[1]])

##              x              y
## [1,] 0.7603140 1.865103e-17
## [2,] 0.8059877 2.998142e-11
## [3,] 0.8516614 1.689587e-06
## [4,] 0.8636612 1.990772e-05
## [5,] 0.8744983 1.599353e-04
## [6,] 0.8754425 1.905200e-04
```

Notes

Organization of the inla-object

Further general information

```
# formula used
result$.args$formula

## y ~ 1 + x
## NULL

# data used
result$.args$data[1:3,]

##           x           y
## 1 0.3457213 1.648559
## 2 0.4242941 1.910131
## 3 0.8127039 2.389109


# log-file including information of INLA approximations
result$logfile
```

Notes

Marginal posterior densities

The marginal posterior densities are stored as a matrices with x - and y -values

```
intercept = data.frame(result$marginals.fixed$(Intercept)`)
x = data.frame(result$marginals.fixed$x)
p1 = ggplot(data = intercept) + geom_point(aes(x,y))
p2 = ggplot(data = x) + geom_point(aes(x,y))
p1+p2
```




Notes

Marginal posterior densities

The rough shape can be interpolated to higher resolution using the `inla.smarginal()` function:

```
smoother_dens = data.frame(inla.smarginal(intercept))
ggplot(data = smoother_dens) + geom_point(aes(x,y))
```



Notes

Marginal posterior densities

Manipulation of the computed posterior marginals is possible through the `inla.*marginal()` functions:

```
# compute the 0.05 quantile
inla.qmarginal(0.05, intercept)

## [1] 0.9511269

# Distribution function
inla.pmarginal(0.975, intercept)

## [1] 0.2742024

# Density function
inla.dmarginal(1, intercept)

## [1] 15.53033

# Generate realizations
inla.rmarginal(4, intercept)

## [1] 0.9730758 0.9857803 0.9878979 0.9635838
```

Notes

Other `inla.*marginal()` functions.

Function Name	Usage
<code>inla.dmarginal(x, marginal, ...)</code>	Density at a vector of evaluation points x
<code>inla.pmarginal(q, marginal, ...)</code>	Distribution function at a vector of quantiles q
<code>inla.qmarginal(p, marginal, ...)</code>	Quantile function at a vector of probabilities p .
<code>inla.rmarginal(n, marginal)</code>	Generate n random deviates
<code>inla.hpdmarginal(p, marginal, ...)</code>	Compute the highest posterior density interval at level p
<code>inla.emarginal(fun, marginal, ...)</code>	Compute the expected value of the marginal assuming the transformation given by <code>fun</code>
<code>inla.mmarginal(marginal)</code>	Compute the mode
<code>inla.smarginal(marginal, ...)</code>	Smoothed density in form of a list of length two. The first entry contains the x-values, the second entry includes the interpolated y-values
<code>inla.tmarginal(fun, marginal, ...)</code>	Transform the marginal using the function <code>fun</code> .
<code>inla.zmarginal(marginal)</code>	Summary statistics for the marginal

Notes

Add random effects

Notes

Getting INLA
○○○○

Implementing the INLA algorithm
○○○○○

How to use INLA
○○○○○○

Simple example
○○○○○○○○○○

Add random effects
●●○○○

Prediction
○○○○○○

Smoothing binary time series
○○○○○○

Disease Mapping
○○○○○○○

Changing the prior
○○○

Repeated Poisson counts
○○○○○○○○○○○○○○

Control statements
○○○

Add random effects

```
f(name, model="...", hyper=...,  
  constr=FALSE, cyclic=FALSE, ...)
```

- **name** – the index of the effect (each f-function needs its own!)
- **model** – the type of latent model. E.g. iid, rw2, ar1, besag, and so on
- **hyper** – specify the prior on the hyperparameters
- **constr** – sum-to-zero constraint?
- **cyclic** – are you cyclic?
- ...

Notes

Getting INLA
○○○○

Implementing the INLA algorithm
○○○○○

How to use INLA
○○○○○○

Simple example
○○○○○○○○○○

Add random effects
●●○○○

Prediction
○○○○○○

Smoothing binary time series
○○○○○○

Disease Mapping
○○○○○○○

Changing the prior
○○○

Repeated Poisson counts
○○○○○○○○○○○○○○

Control statements
○○○

Example: Add random effect

Add an AR(1) random effect to the linear predictor.

- **Stage 1:**
$$y_i|\eta_i \sim \mathcal{N}(\eta_i, \sigma^2)$$
- **Stage 2:** Covariates and AR(1) component connected to likelihood by
$$\eta_i = \beta_0 + \beta_1 x_i + a_i$$
- **Stage 3:**
 - σ^2 : variance of observation noise
 - ρ : dependence in AR(1) process
 - σ^2 : variance of the innovations in AR(1) process

Notes

Getting INLA
○○○○

Implementing the INLA algorithm
○○○○○

How to use INLA
○○○○○○

Simple example
○○○○○○○○○○

Add random effects
●●○○○

Prediction
○○○○○○

Smoothing binary time series
○○○○○○

Disease Mapping
○○○○○○○

Changing the prior
○○○

Repeated Poisson counts
○○○○○○○○○○○○○○

Control statements
○○○

Example: Add random effect


```
# Generate AR(1) sequence  
set.seed(580258)  
t = 1:100  
rho = 0.8  
sd_ar1 = 0.1  
ar = rep(0,100)  
for(i in 2:100)  
  ar[i] = rho * ar[i-1] + rnorm(n = 1, sd = sd_ar1)  
# Generate data with AR(1) component  
x = runif(100)  
y = 1 + 2*x + ar + rnorm(n = 100, sd = 0.2)  
  
# Run inla  
formula = y ~ 1 + x + f(t, model="ar1")  
  
result = inla(formula,  
  data = data.frame(x = x, y = y, t = t),  
  family = "gaussian")
```

Notes

Example

Estimates of the random effect

```
result$summary.random$t %>% ggplot() +
  geom_line(aes(ID, mean)) +
  geom_ribbon(aes(ID, ymin = `0.025quant`, ymax = `0.975quant`), alpha = 0.3) +
  geom_point(data = data.frame(t = t , ar = ar), aes(t,ar), color = "red")
```




Notes

Example

Estimates of the hyperparameters

```
# rho
p1 = ggplot() + geom_line(data = data.frame(result$marginals.hyperpar$`Rho for t`),
  aes(x,y)) +
  geom_vline(xintercept = rho) + ggtitle("rho")
# sd of the ar1 effect
prec = result$marginals.hyperpar$`Precision for t`
sd = inla.tmarginal(function(x) 1/sqrt(x), prec)
p2 = ggplot() + geom_line(data = data.frame(sd), aes(x,y)) +
  geom_vline(xintercept =sd_ar1 ) + ggtitle("sd of ar1")
p1+p2
```



Notes

Prediction

Notes

Getting INLA
○○○○

Implementing the INLA algorithm
○○○○○○

How to use INLA
○○○○○○○

Simple example
○○○○○○○○○○○○

Add random effects
○○○○○○

Prediction
●○○○○○

Smoothing binary time series
○○○○○○○

Disease Mapping
○○○○○○○○

Changing the prior
○○○

Repeated Poisson counts
○○○○○○○○○○○○○○○○

Control statements
○○○

The interpretation of NA

R-INLA uses NA differently than other packages

- NA in the **response** means no likelihood contribution, i.e. response is unobserved
- NA in a **fixed effect** means no contribution to the linear predictor, i.e. the covariate is set equal to zero
- NA in a **random effect** $f(\dots)$ means no contribution to the linear predictor

Notes

Getting INLA
○○○○

Implementing the INLA algorithm
○○○○○○

How to use INLA
○○○○○○○

Simple example
○○○○○○○○○○○○

Add random effects
○○○○○○

Prediction
●○○○○○

Smoothing binary time series
○○○○○○○

Disease Mapping
○○○○○○○○

Changing the prior
○○○

Repeated Poisson counts
○○○○○○○○○○○○○○○○

Control statements
○○○

Prediction

The distribution of the linear predictor at an unobserved location can be computed by specifying the value of the covariate x and the desired time index t and set y to NA.

```
# Add one new location
n = 1
x = c(x, runif(n))
t = c(t, 101:(100+n))
y = c(y, rep(NA,n))

# Re-compute
result.pred = inla(formula,
  data = data.frame(x = x, t = t, y = y),
  family="gaussian",
  control.inla = list(int.strategy = "grid"),
  control.compute = list(config = TRUE,
    return.marginals.predictor=TRUE),
  # tell inla to return the marginals for eta!
  control.predictor = list(compute = TRUE))
```

Notes

Getting INLA
○○○○

Implementing the INLA algorithm
○○○○○○

How to use INLA
○○○○○○○

Simple example
○○○○○○○○○○○○

Add random effects
○○○○○○

Prediction
●○○○○○

Smoothing binary time series
○○○○○○○

Disease Mapping
○○○○○○○○

Changing the prior
○○○

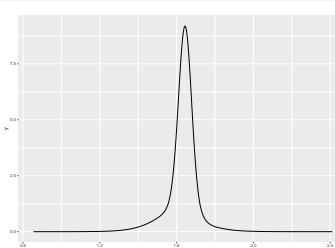
Repeated Poisson counts
○○○○○○○○○○○○○○○○

Control statements
○○○

Prediction

Predicted marginal of the linear predictor η_{101}

```
pred = result.pred$marginals.linear.predictor[[100+n]]
pred = inla.smarginal(pred)
ggplot() +
  geom_line(data = data.frame(pred), aes(x, y))
```



Notes

Prediction

Caution: This is **not** yet the predictive distribution, as the observation noise is missing.

The predictive distribution is

$$\pi(y_{101}|\mathbf{y})$$

what we got is

$$\pi(\eta_{101}|\mathbf{y})$$

Notes

Prediction

One way to add the observation noise to the linear predictor is by sampling from the posterior distribution.

```
n = 1000
x = inla.posterior.sample(n, result.pred)


func = function(...)
{
  eta = Predictor
  eta = eta[101]
  sd = 1/sqrt(theta[1])
  out = rnorm(1, mean = eta, sd =sd)
  return(out)
}

samples = inla.posterior.sample.eval(func, x)[1,]
```

Notes

Prediction

Comparing $\pi(y_{101}|\mathbf{y})$ and $\pi(\eta_{101}|\mathbf{y})$



Notes

Getting INLA	Implementing the INLA algorithm	How to use INLA	Simple example	Add random effects	Prediction
oooo	ooooooo	ooooooo	oooooooooooo	oooooo	ooooooo

Smoothing binary time series	Disease Mapping	Changing the prior	Repeated Poisson counts	Control statements
●oooooo	ooooooooo	ooo	oooooooooooooooooooo	ooo

Notes


Smoothing binary time series

Getting INLA	Implementing the INLA algorithm	How to use INLA	Simple example	Add random effects	Prediction
oooo	ooooooo	ooooooo	oooooooooooo	oooooo	ooooooo

Smoothing binary time series	Disease Mapping	Changing the prior	Repeated Poisson counts	Control statements
●oooooo	ooooooooo	ooo	oooooooooooooooooooo	ooo

Example: Smoothing binary time series

The data set **Tokyo** is available in the **INLA** package and consists of the number of days in Tokyo with rainfall above 1 mm in 1983–1984.



Notes

Getting INLA	Implementing the INLA algorithm	How to use INLA	Simple example	Add random effects	Prediction
oooo	ooooooo	ooooooo	oooooooooooo	oooooo	ooooooo

Smoothing binary time series	Disease Mapping	Changing the prior	Repeated Poisson counts	Control statements
●oooooo	ooooooooo	ooo	oooooooooooooooooooo	ooo

Observations

Each observation consists of

- t : Day of year; $t \in \{1, 2, \dots, 366\}$
- n_t : Number of observations for day t in 1983–1984;
 $n_t \in \{1, 2\}$
- y_t : Number of days with rain out of n_t days for day t ;
 $y_t \in \{0, 1, 2\}$

```
data(Tokyo)
head(Tokyo,3)

##   y n time
## 1 0 2    1
## 2 0 2    2
## 3 1 2    3
Tokyo[60,]
```

```
##   y n time
## 60 0 1   60
```

Notes

Getting INLA
○○○○

Implementing the INLA algorithm
○○○○○○

How to use INLA
○○○○○○○

Simple example
○○○○○○○○○○○○

Add random effects
○○○○○○

Prediction
○○○○○○○

Smoothing binary time series
○○○●○○○

Disease Mapping
○○○○○○○○

Changing the prior
○○○

Repeated Poisson counts
○○○○○○○○○○○○○○○○

Control statements
○○○

Hierarchical model

- **Stage 1:** We have binomial responses with known n_t , but unknown probabilities
$$y_t \sim \text{Binomial}(n_t, p_t)$$
- **Stage 2:** A cyclic second order random walk (CRW2) is connected to the likelihood by
$$p_t = \frac{\exp(\eta_t)}{1 + \exp(\eta_t)}$$
 with linear predictor $\eta_t = \text{CRW2}_t$
- **Stage 3:**
 - τ : Scale parameter in CRW2 with prior
$$\pi(\tau) \sim \text{Gamma}(1, 5 \cdot 10^{-5})$$

Notes

Getting INLA
○○○○

Implementing the INLA algorithm
○○○○○○

How to use INLA
○○○○○○○

Simple example
○○○○○○○○○○○○

Add random effects
○○○○○○

Prediction
○○○○○○○

Smoothing binary time series
○○○○●○○

Disease Mapping
○○○○○○○○

Changing the prior
○○○

Repeated Poisson counts
○○○○○○○○○○○○○○○○

Control statements
○○○

INLA implementation

```
# Read data
data(Tokyo)
# Specify linear predictor
formula = y ~ -1 + f(time, model="rw2", cyclic=TRUE)
# Run model
result = inla(formula,
               family = "binomial",
               Ntrials = n,
               data = Tokyo)
```

Notes

Getting INLA
○○○○

Implementing the INLA algorithm
○○○○○○

How to use INLA
○○○○○○○

Simple example
○○○○○○○○○○○○

Add random effects
○○○○○○

Prediction
○○○○○○○

Smoothing binary time series
○○○○○●○

Disease Mapping
○○○○○○○○

Changing the prior
○○○

Repeated Poisson counts
○○○○○○○○○○○○○○○○

Control statements
○○○


Marginal posterior of CRW2

```
ggplot(data = result$summary.random$t) +
  geom_line(aes(ID, mean)) +
  geom_ribbon(aes(ID, ymin = `0.025quant`, ymax = `0.975quant` ),
            alpha = 0.5)
```

Notes

Transform to probability

```
pred = result$summary.fitted.values
pred$ID = 1:dim(Tokyo)[1]
ggplot(data = pred) +
  geom_line(aes(ID, mean)) +
  geom_ribbon(aes(ID, ymin = `0.025quant`, ymax = `0.975quant` ),
    alpha = 0.5)
```



Notes


Disease Mapping

Notes

Example: disease mapping

We observed larynx cancer mortality counts for males in 544 district of Germany from 1986 to 1990 and want to make a model.

- y_i : The count at location i .
- E_i : An offset; expected number of cases in district i .
- c_i : A covariate (level of smoking consumption) at i
- s_i : spatial location i .



Notes

Disease mapping

Assume

$Y_i \mid \eta_i \sim \text{Poisson}(E_i \exp(\eta_i))$

where the log relative risk is decomposed into

$\eta_i = \mu + u_i + v_i$

- μ is the overall level (intercept).
- $v_i \sim \mathcal{N}(0, \tau_v^{-1})$ represents non-spatial overdispersion.
- u_i are random effects with spatial structure.

Notes

A spatially structured effect

To incorporate a spatial structure into a model, the so called **Besag model** is often used.

$$p(\mathbf{u} \mid \kappa_{\mathbf{u}}) \propto \kappa_u^{(n-1)/2} \exp\left(-\frac{\kappa_u}{2} \sum_{i \sim j} (u_i - u_j)^2\right)$$
$$= \kappa_u^{(n-1)/2} \exp\left(-\frac{\kappa_u}{2} \mathbf{u}^T \mathbf{R} \mathbf{u}\right).$$

where R is called structure matrix and defined as

$$R_{ij} = \begin{cases} n_i & i = j \\ -1 & i \sim j \\ 0 & \text{otherwise.} \end{cases}$$

Here, $i \sim j$ denotes that i and j are neighbouring regions.

Notes

What does this mean?

Example: Five counties of the US state Rhode Island

The structure matrix **R** defines the neighborhood structure.




Table 1: Structure matrix **R**

With increasing number of regions **R** will be sparse, which allows to do many computations very efficient.

Notes

Getting INLA

Implementing the INLA algorithm

How to use INLA

Simple example

Add random effects

Prediction

Smoothing binary time series

Disease Mapping

Changing the prior

Repeated Poisson counts

Control statements

INLA code

```
library(spam)
# load the dataset
data(Oral)
# load the file including neighbourhood information
g = system.file("demodata/germany.graph", package="INLA")
# add one column
Oral = cbind(Oral, region = 1:544, region.unstruc= 1:544)
# define formula
formula = Y ~ f(region, model="besag", graph=g) +
              f(region.unstruc, model="iid")
# run the model
result = inla(formula, family="poisson", E=E, data=Oral)
```

Notes

Getting INLA

Implementing the INLA algorithm

How to use INLA

Simple example

Add random effects

Prediction

Smoothing binary time series

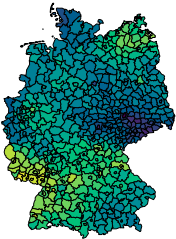
Disease Mapping

Changing the prior

Repeated Poisson counts

Control statements

Median of u on exp-scale



Notes

Getting INLA

Implementing the INLA algorithm

How to use INLA

Simple example

Add random effects

Prediction

Smoothing binary time series

Disease Mapping

Changing the prior

Repeated Poisson counts

Control statements

Other choices for **f**-terms

## [1]	"linear"	"iid"	"mec"	"meb"	"rgeneric"
## [6]	"cgeneric"	"rw1"	"rw2"	"crw2"	"seasonal"
## [11]	"besag"	"besag2"	"bym"	"bym2"	"besagproper"
## [16]	"besagproper2"	"fgn"	"fgn2"	"ar1"	"arlc"
## [21]	"ar"	"ou"	"intslope"	"generic"	"generic0"
## [26]	"generic1"	"generic2"	"generic3"	"spde"	"spde2"
## [31]	"spde3"	"iidid"	"iid2d"	"iid3d"	"iid4d"
## [36]	"iid5d"	"iidkd"	"2diid"	"z"	"rw2d"
## [41]	"rw2diid"	"slm"	"matern2d"	"dmatern"	"copy"
## [46]	"clinear"	"sigm"	"revsigm"	"log1exp"	"logdist"

Notes

Getting INLA	Implementing the INLA algorithm	How to use INLA	Simple example	Add random effects	Prediction	Smoothing binary time series	Disease Mapping	Changing the prior	Repeated Poisson counts	Control statements
0000	000000	0000000	00000000000	000000	0000000	0000000	00000000	000	0000000000000000	000

Notes

Changing the prior

Getting INLA	Implementing the INLA algorithm	How to use INLA	Simple example	Add random effects	Prediction	Smoothing binary time series	Disease Mapping	Changing the prior	Repeated Poisson counts	Control statements
0000	000000	0000000	00000000000	000000	0000000	0000000	00000000	000	0000000000000000	000

Changing the prior: Internal scale

- Hyperparameters are represented internally with more well-behaved transformations, e.g. correlation ρ and precision τ are internally
$$\theta_1 = \log(\tau)$$
$$\theta_2 = \log\left(\frac{1 + \rho}{1 - \rho}\right)$$
- The prior must be set on the parameter in **internal scale**
- Initial values for the mode-search must be set in **internal scale**
- The functions `to.theta()` and `from.theta()` can be used to map back and forth.

Notes

Getting INLA	Implementing the INLA algorithm	How to use INLA	Simple example	Add random effects	Prediction	Smoothing binary time series	Disease Mapping	Changing the prior	Repeated Poisson counts	Control statements
0000	000000	0000000	00000000000	000000	0000000	0000000	00000000	000	0000000000000000	000

Changing the prior: Code

```
hyper = list(prec = list(prior = "loggamma",
                        param = c(1, 0.1),
                        initial = 4,
                        fixed = FALSE))

formula = y ~ f(idx, model = "iid", hyper = hyper) + ...

# For the iid model, default options can be seen with
inla.doc("iid")
```

Notes

Getting INLA	Implementing the INLA algorithm	How to use INLA	Simple example	Add random effects	Prediction	Smoothing binary time series	Disease Mapping	Changing the prior	Repeated Poisson counts	Control statements
oooo	oooooo	ooooooo	oooooooooooo	oooooo	ooooooo	ooooooo	ooooooo	ooo	oooooooooooooooo	ooo

Repeated Poisson counts

Notes

Getting INLA	Implementing the INLA algorithm	How to use INLA	Simple example	Add random effects	Prediction	Smoothing binary time series	Disease Mapping	Changing the prior	Repeated Poisson counts	Control statements
oooo	oooooo	ooooooo	oooooooooooo	oooooo	ooooooo	ooooooo	ooooooo	ooo	oooooooooooooooo	ooo

EPIL example

Seizure counts in a randomised trial of anti-convulsant therapy in epilepsy. From WinBUGS manual.

```
## # A tibble: 6 x 8
##   Ind Rep11 Rep12 Rep13 Rep14   Trt   Base   Age
##   <int> <dbl> <dbl> <dbl> <dbl> <int> <int> <int>
## 1     1     5     3     3     3     0    11    31
## 2     2     3     5     3     3     0    11    30
## 3     3     2     4     0     5     0     6    25
## 4     4     4     4     1     4     0     8    36
## 5     5     7    18     9    21     0    66    22
## 6     6     5     2     8     7     0    27    29
```

Covariates are treatment (0,1), 8-week baseline seizure counts, and age in years.

Notes

Getting INLA	Implementing the INLA algorithm	How to use INLA	Simple example	Add random effects	Prediction	Smoothing binary time series	Disease Mapping	Changing the prior	Repeated Poisson counts	Control statements
oooo	oooooo	ooooooo	oooooooooooo	oooooo	ooooooo	ooooooo	ooooooo	ooo	oooooooooooooooo	ooo

Repeated Poisson counts

$$\begin{aligned} y_{jk} &\sim \text{Poisson}(\mu_{jk}); \quad j = 1, \dots, 59; \quad k = 1, \dots, 4 \\ \log(\mu_{jk}) &= \alpha_0 + \alpha_1 \log(\text{Base}_j/4) + \alpha_2 \text{Trt}_j \\ &\quad + \alpha_3 \text{Trt}_j \log(\text{Base}_j/4) + \alpha_4 \log(\text{Age}_j) \\ &\quad + \alpha_5 V4 + \text{Ind}_j + \beta_{jk} \\ \alpha_i &\sim \mathcal{N}(0, \tau_\alpha) \quad \tau_\alpha \text{ known (0.001)} \\ \text{Ind}_j &\sim \mathcal{N}(0, \tau_{\text{Ind}}) \quad \tau_{\text{Ind}} \sim \text{Gamma}(1, 0.01) \\ \beta_{jk} &\sim \mathcal{N}(0, \tau_\beta) \quad \tau_\beta \sim \text{Gamma}(1, 0.01) \end{aligned}$$

Here, V4 is an indicator variable for the 4th visit.

Notes

Getting INLA
oooo

Implementing the INLA algorithm
oooooo

How to use INLA
ooooooo

Simple example
oooooooooooo

Add random effects
oooooo

Prediction
ooooooo

Smoothing binary time series
ooooooo

Disease Mapping
ooooooo

Changing the prior
ooo

Repeated Poisson counts
ooo●oooooooooooo

Control statements
ooo

Model specification in INLA

The data:

```
##   y Trt Base Age V4 rand Ind
## 1 5   0  11  31  0   1   1
## 2 3   0  11  31  0   2   1
## 3 3   0  11  31  0   3   1
## 4 3   0  11  31  1   4   1
## 5 3   0  11  30  0   5   2
## 6 5   0  11  30  0   6   2
```

The formula:

```
formula = y ~ ClBase4*CTrt + ClAge + CV4 +
  f(Ind, model="iid",
    hyper = list(prec = list(prior = "loggamma",
                             param = c(1,0.01)))) +
  f(rand, model="iid",
    hyper = list(prec = list(prior = "loggamma",
                             param = c(1,0.01))))
```

Run the model:

```
result = inla(formula, family="poisson", data = Epil,
  control.fixed = list(prec.intercept = 0.001,
    prec = 0.001))
```

Notes

Getting INLA
oooo

Implementing the INLA algorithm
oooooo

How to use INLA
ooooooo

Simple example
oooooooooooo

Add random effects
oooooo

Prediction
ooooooo

Smoothing binary time series
ooooooo

Disease Mapping
ooooooo

Changing the prior
ooo

Repeated Poisson counts
ooo●oooooooooooo

Control statements
ooo

Comparing results with MCMC

- When comparing the results of R-INLA with MCMC, it is important to use the **same model**.

That means, same data, same priors, same constraints on parameters, intercept included or not,

Notes

Getting INLA
oooo

Implementing the INLA algorithm
oooooo

How to use INLA
ooooooo

Simple example
oooooooooooo

Add random effects
oooooo

Prediction
ooooooo

Smoothing binary time series
ooooooo

Disease Mapping
ooooooo

Changing the prior
ooo

Repeated Poisson counts
ooo●oooooooooooo

Control statements
ooo

Comparing results with MCMC

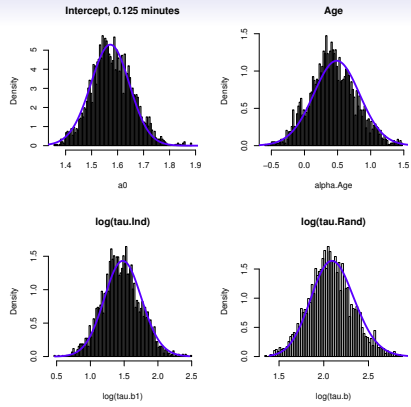
- When comparing the results of R-INLA with MCMC, it is important to use the **same model**.

That means, same data, same priors, same constraints on parameters, intercept included or not,

- Here we have compared the results with those obtained using 'JAGS via the **rjags** package

Notes

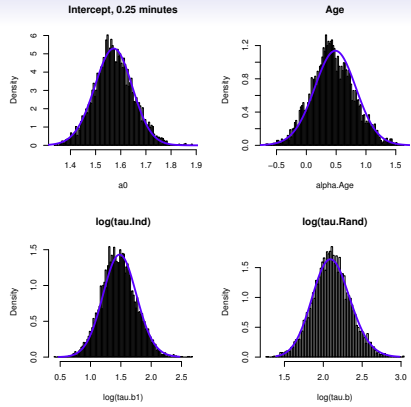
Getting INLA	Implementing the INLA algorithm	How to use INLA	Simple example	Add random effects	Prediction
oooo	oooooooo	ooooooooo	ooooooooooooo	ooooooo	ooooooooo



Running time of INLA < 0.5 seconds

Notes

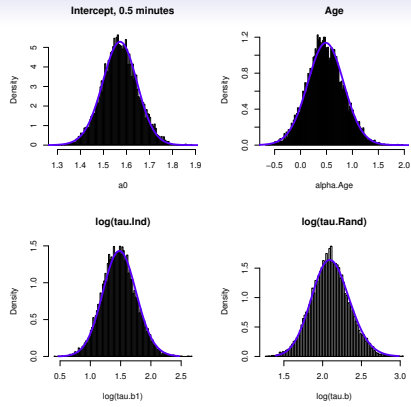
Getting INLA	Implementing the INLA algorithm	How to use INLA	Simple example	Add random effects	Prediction
oooo	oooooooo	ooooooooo	ooooooooooooo	ooooooo	ooooooooo



Running time of INLA < 0.5 seconds

Notes

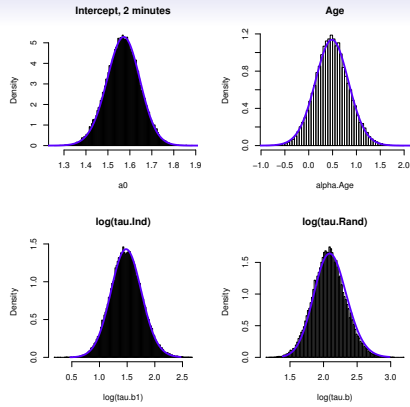
Getting INLA	Implementing the INLA algorithm	How to use INLA	Simple example	Add random effects	Prediction
oooo	oooooooo	ooooooooo	ooooooooooooo	ooooooo	ooooooooo



Running time of INLA < 0.5 seconds

Notes

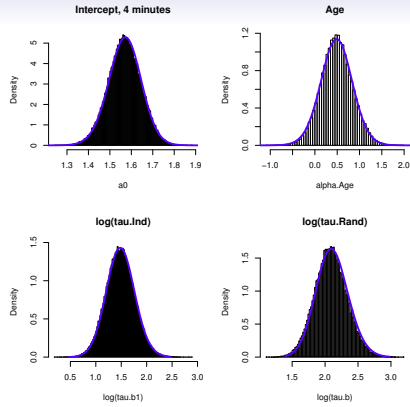
Getting INLA	Implementing the INLA algorithm	How to use INLA	Simple example	Add random effects	Prediction
oooo	oooooooo	oooooooo	oooooooooooo	oooooo	oooooooo



Running time of INLA < 0.5 seconds

Notes

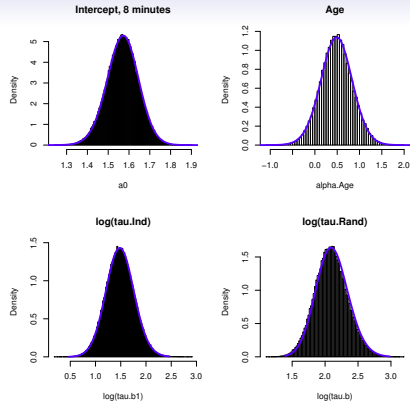
Getting INLA	Implementing the INLA algorithm	How to use INLA	Simple example	Add random effects	Prediction
oooo	oooooooo	oooooooo	oooooooooooo	oooooo	oooooooo



Running time of INLA < 0.5 seconds

Notes

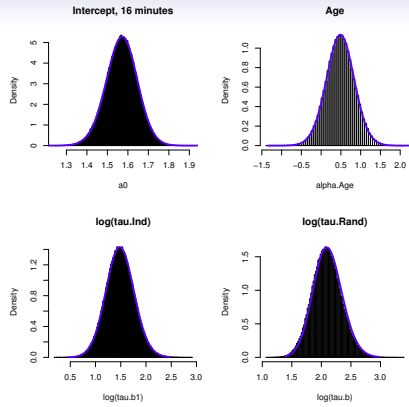
Getting INLA	Implementing the INLA algorithm	How to use INLA	Simple example	Add random effects	Prediction
oooo	oooooooo	oooooooo	oooooooooooo	oooooo	oooooooo



Running time of INLA < 0.5 seconds

Notes

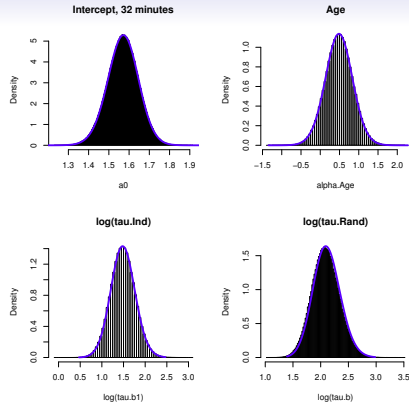
Getting INLA	Implementing the INLA algorithm	How to use INLA	Simple example	Add random effects	Prediction
oooo	oooooooo	ooooooooo	oooooooooooo	oooooo	ooooooo



Running time of INLA < 0.5 seconds

Notes

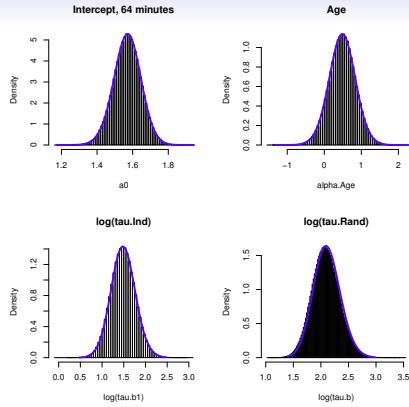
Getting INLA	Implementing the INLA algorithm	How to use INLA	Simple example	Add random effects	Prediction
oooo	oooooooo	ooooooooo	oooooooooooo	oooooo	ooooooo



Running time of INLA < 0.5 seconds

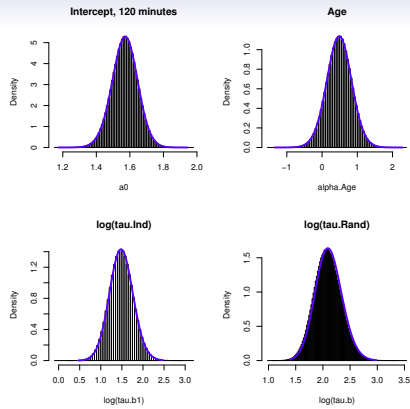
Notes

Getting INLA	Implementing the INLA algorithm	How to use INLA	Simple example	Add random effects	Prediction
oooo	oooooooo	ooooooooo	oooooooooooo	oooooo	ooooooo



Running time of INLA < 0.5 seconds

Notes



Running time of INLA < 0.5 seconds

Notes

Control statements

Notes

Control statements

`control.xxx` statements control computations

- `control.fixed`
 - `prec`: Default precision for all fixed effects except the intercept.
 - `prec.intercept`: Precision for intercept (Default: 0.0)
- `control.predictor`
 - `compute`: Compute posterior marginals of linear predictors
- `control.compute`
 - `dic, mlik, cpo`: Compute measures of fit
 - `config`: Save internal GMRF approximations? (needed to use `inla.posterior.sample()`)
- `control.inla`
`strategy` and `int.strategy` contain useful advanced features
- There are various others as well; see help.

Notes

Getting INLA	Implementing the INLA algorithm	How to use INLA	Simple example	Add random effects	Prediction	Smoothing binary time series	Disease Mapping	Changing the prior	Repeated Poisson counts	Control statements
oooo	oooooooo	ooooooooo	ooooooooooooo	ooooooo	ooooooooo	ooooooooo	oooooooooo	ooo	oooooooooooooooooooo	oo●

Thank you for your attention!

If you have any doubts or questions, please write :
sara.martino@math.ntnu.no



Notes

Notes

Notes
