

# Spatial modeling with INLA and inlabru

University of Zurich, March, 2022

Instructor: Sara Martino

Department of Mathematical Science (NTNU)



# NTNU

Norwegian University of  
Science and Technology

## Model choice and model assessment/validation

Advanced features

Feature:Replicate

Feature:Group

Feature:Multiple likelihoods

Feature: copy

Model choice and model assessment/validation	Advanced features	Feature:Replicate	Feature:Group	Feature
oooooooooooooooo	oo	oooooooo	oo	oooo

# Model choice and model assessment/validation

# Introduction

We have now seen some fancy modelling approaches

# Introduction

We have now seen some fancy modelling approaches

How can we assess the models and choose between them?

- Rather underdeveloped in statistical literature; Many suggestions; no clear “yes, this is how it should be done”

## Model choice and assessment

- **Model assessment** is the art and science of evaluating how well a model and/or estimate agrees with observed reality, and of how useful it for specific purposes
  - Simple models -summary characteristics
  - Complex models - assessing variability in space
  - All models - prediction ability; calibrated uncertainty
- **Model choice** - which covariate and random effects to include
- **Model comparison** - which model is “better?”

## Model choice

INLA can compute the following quantities:

- Marginal likelihood  $\Rightarrow$  Bayes factors
- Deviance information criterion (DIC)
- Widely applicable information criterion (WAIC)



## General advice

- We have little experience with practical usage of them for complex spatial models
- It is not clear what they actually mean in the context of the models we look at here
- Advice: use them cautiously
- Less adventurous if you are comparing models with only different numbers of covariates - and “the rest” is the same:
  - Use the same mesh in the models you compare (do not treat the mesh resolution as a model choice!)

## Marginal likelihood

```
result = inla(...,  
              control.compute=list(mlik=TRUE))  
  
result = bru(...,options = list(control.compute =  
                                list(mlik = TRUE)))
```

- Calculates  $\log(\pi(\mathbf{y}))$
- Can calculate Bayes factors through differences in value
- **NB:** Problematic for intrinsic models

## Deviance information criterion

```
result = inla(...,  
              control.compute=list(dic=TRUE))  
  
result = bru(...,options = list(control.compute =  
                                list(dic = TRUE)))
```

DIC is a measure of complexity and fit. It is used to compare complex hierarchical models and is defined as:

$$\text{DIC} = \overline{D} + p_D$$

where  $\overline{D}$  is the posterior mean of the deviance and  $p_D$  is the effective number of parameters. Smaller values of the DIC indicate a better trade-off between complexity and fit of the model.

# Widely applicable information criterion (WAIC)

```
result = inla(...,  
              control.compute=list(waic=TRUE))  
  
result = bru(...,options = list(control.compute =  
                                list(waic = TRUE)))
```

- WAIC is like DIC just newer, and perhaps better
- See “*Understanding predictive information criteria for Bayesian models*” (2013) by Andrew Gelman, Jessica Hwang, and Aki Vehtari

## Model assessment with cross-validated scores

Posterior predictive distributions can be used for model assessment and model selection.

Full cross-validation or out-of-sample validation is expensive.

R-INLA provides two leave-one-out crossvalidation quantities:

- **Conditional predictive ordinate**
- **Probability integral transform**

## Conditional predictive ordinate

```
result = inla(...,  
              control.compute=list(cpo=TRUE))  
  
result = bru(...,options = list(control.compute =  
                                list(cpo = TRUE)))
```

- Measures fit through the predictive density  $\pi(y_i^{obs} | \mathbf{y}_{-i})$
- Basically, Bayesian hold-one out
- Easy to compute in the INLA-approach
- Possible failure (`$cpo$failure`)
- See *Posterior and Cross-validatory Predictive Checks: A Comparison of MCMC and INLA* (2009) by Held, Schrödle and Rue

## Proper scoring rule based on CPO

A predictive score is proper if its expected value is minimised under the true distribution.

- The **log-CPO-score**

$$\log\text{CPO} = - \sum_{i=1}^n \log(\text{CPO}_i) = - \sum_{i=1}^n \log[p(y_i^{\text{obs}} | y_j^{\text{obs}}, j \neq i)]$$

is a strictly proper scoring rule.

- The logCPO score encourages appropriate prediction uncertainty; bias, overconfidence, and underconfidence all increase the score.
- 2logCPO is similar in scale to DIC and WAIC but has a clear cross validation prediction interpretation.

## Pairwise observation CPO scores

- The aggregated logCPO score hides information
- Model comparison for predictions is a pairwise comparison problem for each individual observation!
- Compute the collection of pairwise logCPO differences for two models
- Inspect the empirical score difference distribution; is it consistently positive/negative?
- Inspect the spatial pattern of the score differences



## Probability integral transform

```
result = inla(...,
               control.compute=list(pit=TRUE))

result = bru(...,options = list(control.compute =
                                list(pit = TRUE)))
```

- Given by

$$\text{Prob}(Y_i \leq y_i^{obs} \mid \mathbf{y}_{-i})$$

## PIT: assessing prediction bias, scale and shape

- A direct consequence of the PIT definition is that under the true model, each  $PIT_i$  value is a sample of a uniform distribution on  $[0, 1]$ .
- The usual plotting method for PIT is a histogram.
- For models with too small predictive variance, the histogram tends to increase toward 0 and 1.
- For models with too large predictive variance, the histogram tends to have peak in the middle.
- For incorrectly skewed predictions, the PIT histogram will tend to be skewed.
- Unfortunately, that doesn't necessarily imply that overfitting and oversmoothing can be detected and/or correctly diagnosed.

## Advanced features

## Useful features

There are several features that can be used to extend the standard models in **R-INLA** (and **inlabru**)

- Replicate
- Group
- Copy
- Multiple likelihoods
- Generic precision matrices (**rgeneric**)

### Main goals

- know about the features
- be exposed to the ideas

Model choice and model assessment/validation  
oooooooooooooooo

Advanced features  
oo

**Feature:Replicate**  
●oooooooo

Feature:Group  
oo

Feature:  
oooo

Feature:Replicate

## Feature: replicate

**replicate** generates iid replicates from the same  $f()$ -model with the same hyperparameters.

If  $\mathbf{x} \mid \theta \sim \text{AR}(1)$ , then `nrep=3`, makes

$$\mathbf{x} = (\mathbf{x}_1, \mathbf{x}_2, \mathbf{x}_3)$$

with mutually independent  $\mathbf{x}_i$ 's from  $\text{AR}(1)$  with the same  $\theta$

```
f(..., replicate = r [, nrep = nr ])
```

where replicate are integers 1, 2, ..., etc

## Example

$$y_i^1 \sim \text{Poisson}(\lambda_i^1), \quad i = 1, \dots, n_1$$

$$y_i^2 \sim \text{Poisson}(\lambda_i^2), \quad i = 1, \dots, n_2$$

$$\log(\lambda_i^1) = \mu_1 + u_i^1$$

$$\log(\lambda_i^2) = \mu_2 + u_i^2$$

and  $\mathbf{u}^1$  and  $\mathbf{u}^2$  are two replicates of the same AR1 model (they share the same parameters)


## Example : simulate data

```
# Simulate data - 2 groups with same AR1 param
n = 100
rho <- 0.8
mu = c(1,-1)
x1 = arima.sim(n=n, model=list(ar=c(rho))) + mu[1]
x2 = arima.sim(n=n, model=list(ar=c(rho))) + mu[2]
# generate Poisson observations
y1 = rpois(n, lambda = exp(x1))
y2 = rpois(n, lambda = exp(x2))

df_groups <- data.frame(y = c(y1, y2),
                        t = rep(1:n, 2),
                        repl = rep(1:2, each = n),
                        int = rep(0:1, each = n))
```





## Example : simulate data




## Example: fit the model

```
cmp <- y ~ -1 + int(int, model = "factor_full") +  
  myar1(t, model = "ar1", replicate = repl)  
fit <- bru(cmp, family = "poisson", data = df_groups)
```

## Example: Results - Latent field



## Example: Results - Hyperparameters



Model choice and model assessment/validation  
oooooooooooooooo

Advanced features  
oo

Feature:Replicate  
oooooooo

**Feature:Group**  
●o

Feature:  
oooo

Feature:Group

## Feature: group

- Similar concept as replicate, but with a dependence structure on the replicates. E.g. ~rw1, rw2, ar1, exchangeable
- Implemented as a Kronecker product (often space and time)
- It's possible to use both replicate and group! This will be replications of the grouped model
- Usage

```
f(..., group = g[, ngroup = ng])
```

where replicate are integers 1, 2, ..., etc

Feature:Multiple likelihoods

## Feature:Multiple likelihood

There is no constraint in INLA that the type of likelihood must be the same for all observations. In fact, every observation could have its own likelihood.

- Coregionalization model
- Marked point process
- Joint models of various kinds



## Example: Simulate data

We fit a simple model where we imagine that some data come from a Gaussian and some from a Poisson likelihood:

## Example: Fit the model

```
cmp = ~ Intercept_1(1) + Intercept_2(1) +  
      x1(x1, model = "linear") + x2(x2, model = "linear")  
  
lik1 = like(formula = y1~Intercept_1 + x1,  
            family = "gaussian",  
            exclude = c("Intercept_2","x2"),  
            data = d1)  
  
lik2 = like(formula = y2~Intercept_2 + x2,  
            family = "poisson",  
            exclude = c("Intercept_1","x1"),  
            data = d2)  
  
fit = bru(cmp, lik1,lik2)
```

Model choice and model assessment/validation  
oooooooooooooooo

Advanced features  
oo

Feature:Replicate  
oooooooo

Feature:Group  
oo

Feature:  
oooo

Feature: copy

## Feature: copy

Allows different elements of the same `f(...)` to be in the the same linear predictor.

Without copy we can not (directly) specify the model

$$\eta_i = u_i + u_{i+1} + \dots$$

Sometimes this is necessary

## Feature: copy

The linear predictor

$$\eta_i = u_i + u_{i+1} + \dots$$

can be coded as

```
formula = y ~ f(i, model = "iid")  
          + f(i.plus, copy="i") + ...
```

- The copy-feature, creates internally an additional sub-model which is  $\epsilon$ -close to the target
- Many copies allowed, and copies of copies

## Feature: copy

It is also possible to include scaled copies

$$\eta_i = u_i + \beta u_{i+1} + \dots$$

```
formula = y ~ f(i, model="iid") +  
            f(i.plus, copy="i",  
              hyper = list(beta=list(fixed=FALSE)))  
+ ...
```

This introduces another hyperparameter in the model ( which is fixed to 1 by default).

## Feature: copy

$$y_i^1 \sim \mathcal{N}(\mu_i, \tau)$$

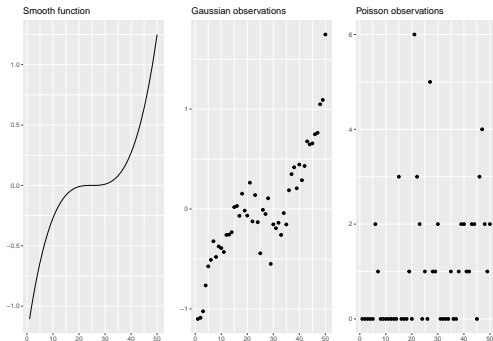
$$\mu_i = f(i)$$

$$y_j^2 \sim \text{Poisson}(\lambda_j), \quad j = 1, \dots, 50$$

$$\log(\lambda_j) = f(j)$$

## Example : simulate data

```
n = 50  
idx = 1:n  
x = idx  
func = 10 * ((idx-n/2)/n)^3  
  
y1 = rnorm(50, mean = func, sd = 0.2)  
y2 = rpois(50, lambda = exp(func))
```





## Example : fit the model

```
df1 = data.frame(y1 = y1, idx1 = 1:n)
df2 = data.frame(y2 = y2, idx2 = 1:n)

cmp = ~ -1 +
  field(idx1, model = "rw1") +
  field_copy(idx2, copy = "field")

lik1 = like(formula = y1~ field,
            family = "gaussian",
            exclude = c("field_copy"),
            data = df1)

lik2 = like(formula = y2~ field_copy,
            family = "poisson",
            exclude = c("field"),
            data = df2)

fit = bru(cmp,
          lik1,
          lik2)
```

## Example: Results

