

Notes

Spatial modeling with INLA and inlabru

University of Zurich, March, 2022

Instructor: Sara Martino

Department of Mathematical Science (NTNU)



Norwegian University of
Science and Technology

Notes

Spatial models

Gaussian Random Field models

The SPDE approach

Fitting spatial models

Space-time Modeling

Prediction in space

Notes

What have we learned

Notes

INLA in a nutshell

- many data sets these days are complex, resulting in complex models, e.g. with complex dependence structures (spatial, temporal, etc..)
 - usually Markov chain Monte Carlo (MCMC) methods have been used to fit these models
 - (realistically) complex models result in very long running times
 - often impossible (or unrealistic) to fit
 - INLA (Integrated nested Laplace approximation) is an alternative to MCMC
 - much, much **faster**
 - suitable for a specific (but very large!) class of models

Notes

INLA in a nutshell

Notes

Three main ingredients in INLA

- Gaussian Markov random fields
 - Latent Gaussian models
 - Laplace approximations

which together (with a few other things) give a very nice tool for Bayesian inference.

- quick
 - accurate

Models amenable to INLA: Latent Gaussian Models

$$\begin{aligned} \mathbf{y}|\mathbf{x}, \theta &\sim \prod_i \pi(y_i|x_i, \theta) && \text{Likelihood} \\ \mathbf{x}|\theta &\sim \exp\left(-\frac{1}{2}\mathbf{x}^T \mathbf{Q}(\theta) \mathbf{x}^T\right) && \text{Latent field (GMRF)} \\ \theta &\sim \pi(\theta) && \text{Hyperparameter} \end{aligned}$$

- Each data point depends on only one of the elements in the latent Gaussian field \mathbf{x} .
- The size of the hyperparameter vector θ is small (say < 15)
- The latent field \mathbf{x} , can be large but it is endowed with some conditional independence (Markov) properties so that the precision matrix $Q(\theta)$ is sparse.
- The linear predictor depends **linearly** on the unknown smooth function of covariates.
- The inferential interest lies in the univariate posterior marginals $\pi(x_i|\mathbf{y})$ and $\pi(\theta_j|\mathbf{y})$ rather than in the joint posterior $\pi(\mathbf{x}, \theta|\mathbf{y})$

Notes

INLA allows us to compute in a **fast** and **efficient** way an **accurate** approximation to the posterior marginals for the hyperparameters

$$\tilde{\pi}(\theta_j|\mathbf{y})$$

and for the latent field

$$\tilde{\pi}(x_i|\mathbf{y})$$

All of this is implemented in the **INLA** library in **R**.

Yesterday we saw how to implement simple models with **INLA**

Today we are going to look at how to implement continuous spatial models.

Such model are **a lot** easier to implement in **inlabru!!!!**

Notes

inlabru in a nutshell

inlabru is a friendlier version of **R-INLA**

- it makes **INLA** more accessible to the user
- makes complex features and predictions (especially for spatial data) a lot easier
- is a softer-wrapper around **INLA**
- in some cases it allows to release some of the constraints of **R-INLA** (the linearity of the predictor)

Notes

inlabru

- Installation:

- There is a CRAN version

```
install.packages("inlabru")
```

- You can also install the development version of inlabru from GitHub (recommended)

```
# install.packages("remotes")
remotes::install_github("inlabru-org/inlabru",
  ref="devel")
```

- Documentation

- Web site: <https://sites.google.com/inlabru.org/inlabru>
- Github: <https://github.com/inlabru-org/inlabru>

Notes

Notes

Spatial models

Types of Spatial Data

Notes

We can distinguish three types of spatial data

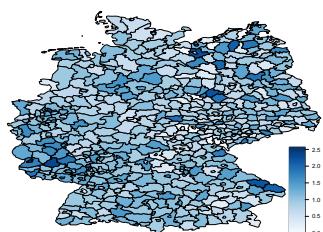
- Discrete space
 - data on a spatial grid
- Continuous space:
 - geostatistical data
 - spatial point data

Notes

Discrete Counts:

Data on a spatial grid

- examples: number of individuals in a region, average rainfall in a province
 - (originally geostatistical or point data; gridded for practical reasons)



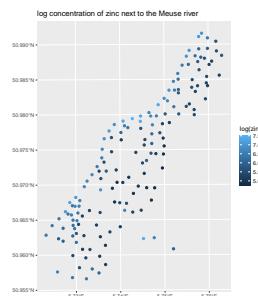
Observed response(s):

- Measurement over each grid cell (e.g. number of individuals in cell; rainfall in province)

Notes

Continuous Space: Geostatistics

- phenomenon that is continuous in space
 - examples: nutrient levels in soil, salinity in the sea
 - measurements at a given set of locations that are determined by surveyor



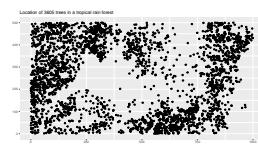
Observed response(s):

- measurement(s) taken at given locations

Notes

Continuous Space: Point Process

- locations of objects (individuals) in space (typically 2D)
 - examples: locations of trees in a forest, groups of animals



Observed response(s):

- x, y coordinates of points (individuals/groups)
 - maybe also properties of individuals/groups (“marks”)

Notes

Notes

Gaussian Random Field models

Gaussian Random fields

Definition: A random function $u(x) : R^d \rightarrow R$ is a Gaussian random field if for any finite collection of locations, (x_1, \dots, x_n) , $x_i \in R^d$, the joint distribution of $\mathbf{u} = (u(x_1), \dots, u(x_n))$ is $\mathbf{u} \sim N(0, \Sigma)$, and

$$E(u(x)) = 0, \\ Cov(u(x), u(x')) = R(x, x'), \quad \Sigma_{ii} = R(x_i, x_j)$$

for some expectation function $\mu(\cdot)$ and positive definite covariance function $R(\cdot, \cdot)$. Σ is the covariance matrix for the specific location collection.

Matern covariance function

The Matern covariance function is a popular model in spatial problems:

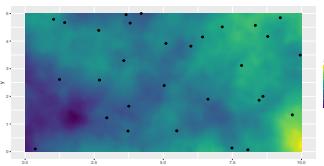
$$c_\nu(d; \sigma, \rho) = \sigma^2 \frac{2^{1-\nu}}{\Gamma(\nu)} \left(\sqrt{8\nu} \frac{d}{\rho} \right)^\nu K_\nu \left(\sqrt{8\nu} \frac{d}{\rho} \right)$$

* ρ and σ are the range and the marginal standard deviation

- K_ν modified Bessel function of the second kind, order ν .

Example

Notes



Gaussian Random field

Notes

GRF are a very popular model

- Flexible and easy to use
 - Can be part of the latent Gaussian field in a LGM

However:

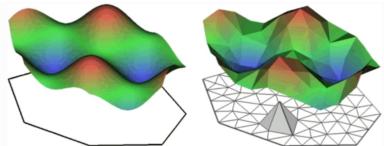
- **computationally inefficient** (the precision matrix is dense)
 - not flexible enough (complicated boundary, barrier, . . .)

The SPDE approach

Notes

The SPDE approach

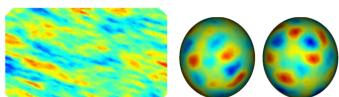
- Matern fields can be seen as solution to a PDE
 - Using finite element methods such solution can be represented using a GRMF



Notes

Advantages of the SPDE approach

- Computationally fast
 - Allows for flexible modeling
 - non-stationary models (anisotropy)
 - models on a sphere
 - non separable models



All these models (and my more) can be fitted with **R-INLA** and **inlabru**.

Notes

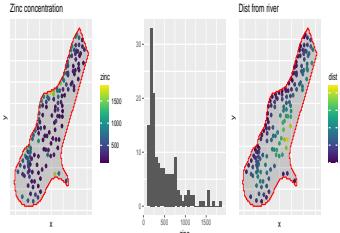
Learning about the SPDE approach

- F. Lindgren, H. Rue, and J. Lindström. *An explicit link between Gaussian fields and Gaussian Markov random fields: The SPDE approach (with discussion)*. In: Journal of the Royal Statistical Society, Series B 73.4 (2011), pp. 423–498.
 - H. Bakka, H. Rue, G. A. Fuglstad, A. Riebler, D. Bolin, J. Illian, E. Krainski, D. Simpson, and F. Lindgren. *Spatial modelling with R-INLA: A review*. In: WIREs Computational Statistics 10:e1443.6 (2018). (Invited extended review). DOI: 10.1002/wics.1443.
 - E. T. Krainski, V. Gómez-Rubio, H. Bakka, A. Lenzi, D. Castro-Camilio, D. Simpson, F. Lindgren, and H. Rue. *Advanced Spatial Modeling with Stochastic Partial Differential Equations using R and INLA*. Github version www.r-inla.org/spde-book. CRC press, Dec. 20

Notes

Notes

Fitting spatial models



Example: Meuse Data

Notes

Measures of zinc concentration.

The model

Notes

$$\log(Y(s)) \sim \mathcal{N}(\eta(s), \sigma_y^2)$$

$$\eta(s) = \alpha + \beta x(s) + u(s)$$

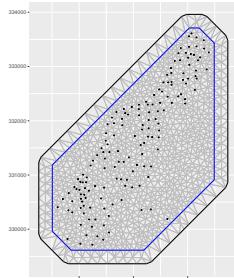
where

- $Y(s)$ is the measure of zinc in location s
 - α a common intercept
 - β a model parameter
 - $x(s)$ distance from the river at location s
 - $u(s)$ the Matern Gaussian field

Step 1: Define the SPDE representation

Notes

- The mesh
 - The SPDE model



Define the SPDE representation: The mesh

1. Define the mesh

```
mesh <- inla.mesh.2d(loc.domain = cbind(meuse$x, meuse$y),
                      max.edge = c(150, 500),
                      offset = c(100, 250) )
```

Notes

- ## Define the SPDE representation: The mesh

 - All random field models need to be discretised for practical calculations.
 - The SPDE models were developed to provide a consistent model definition across a range of discretisations.
 - We use finite element methods with local, piecewise linear basis functions defined on a triangulation of a region of space containing the domain of interest.
 - Deviation from stationarity is generated near the boundary of the region.
 - The choice of region and choice of triangulation affects the numerical accuracy.

Two separate issues:

- Continuous space with bounded domain: Boundary effect
 - Discretised model: Numerical accuracy

Sometimes the boundary effect may be desirable.

Define the SPDE representation: The mesh

Notes

- Too fine meshes → heavy computation
 - Too coarse mesh → not accurate enough

Some guidelines

Notes

- Create triangulation meshes with `inla.mesh.2d()`
 - Move undesired boundary effects away from the domain of interest by extending to a smooth external boundary (`inla.nonconvex.hull(loc, convex)`, `convex ≥ correlation range`)
 - Use a coarser resolution in the extension to reduce computational cost (`max.edge=c(inner, outer)`)
 - Use a fine resolution (subject to available computational resources) for the domain of interest (inner correlation range) and filter out small input point clusters ($0 < \text{cutoff} < \text{inner}$)
 - Coastlines and similar can be added to the domain specification in `inla.mesh.2d()`

Define the SPDE representation: The mesh

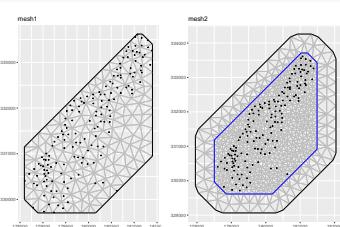
Notes

```

mesh1 = inla.mesh.2d(loc.domain = cbind(meuse$x, meuse$y),
                     max.edge = 350,
                     offset = 10)

mesh2 = inla.mesh.2d(loc.domain = cbind(meuse$x, meuse$y),
                     max.edge = c(150, 500),
                     cutoff = 100,
                     offset = c(100, 550) )

```



Define the SPDE representation: The SPDE model

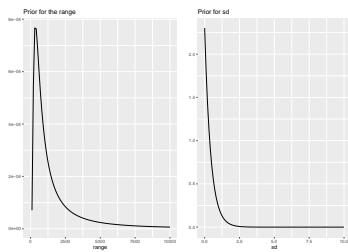
```
meuse.spde <- inla.spde2.pcmatern(mesh = mesh,
                                      prior.sigma = c(1, 0.1),
                                      prior.range = c(1000, 0.5))
```

PC-priors for the range ρ and the standard deviarion σ

- Define the prior for the range `prior.range = (range0,Prange)` Prob($\rho < \rho_0$) = p_ρ
- Define the prior for the range `prior.sigma = (sigma0,Psigma)` Prob($\sigma > \sigma_0$) = p_σ

Notes

Prior for range ρ and sd σ



Notes

Run the model `inlabru`

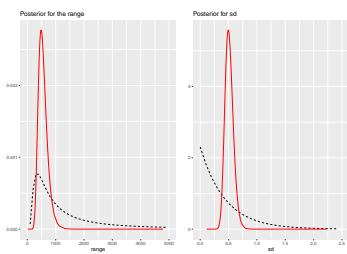
```
# create a spatial object
coordinates(meuse) = c("x", "y")
# covariate values
dist_SPDE = SpatialPixelsDataFrame(data$dist_raster[,c(1,2)],
                                     data = data.frame(dist = data$dist_raster[,3]))
# model components
cmp = - Intercept(1) + dist(dist_SPDE, model = "linear") +
    spde(coordinates, model = meuse.spde)
# define likelihood
lik = like(formula = Y ~ Intercept + dist + spde,
           family = "gaussian",
           data = meuse)
#fit the model
fit <- brm(cmp, lik)
# define prediction area
pix <- pixels(mesh, nx = 200, ny = 200, mask = boundary)
# generate predictions
pred = predict(fit, pix, - data.frame(
    spde = spde,
    logscale = Intercept + dist + spde,
    naturalscale = exp(Intercept + dist + spde)))
```

Notes

Notes!

- The data are a spatial object!
 - For prediction, the covariates are stored in a `SpatialPixelsDataFrame` and need to cover all the mesh nodes

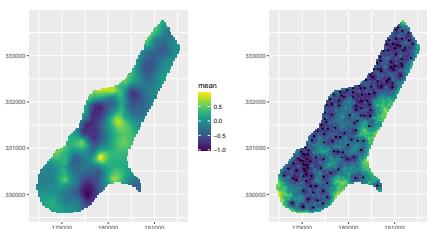
Notes



Notes

Estimation: Posterior for the hyperparameters

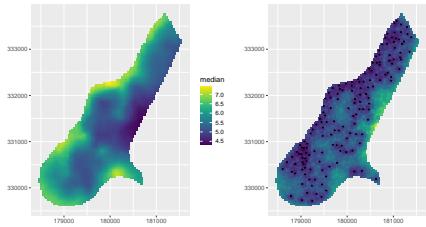
Notes



Notes

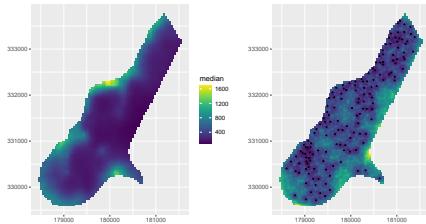
Predictions: The log concentrations

Notes



Predictions: The concentrations

Notes



Same in plain INLA (1)

Notes

```

A.mouse <- inla.spde.make.A(mesh = mesh, loc = coordinates(meuse))
s.index <- inla.spde.make.index(name = "spatial.field",
s.spde <- meuse.spde$spde

#Create data structure
meuse.stack <- inla.stack(data = list(zinc = meuse$zinc),
A = list(A.mouse, 1),
effects = list(c(s.index, list(Intercept = 1)),
list(dist = meuse$dist)),
tag = "meuse.data")

data(meuse.grid)
coordinates(meuse.grid) = -x+y
gridded(meuse.grid) = TRUE

#Create data structure for prediction
A.pred <- inla.spde.make.A(mesh = mesh, loc = coordinates(meuse.grid))
meuse.stack <- inla.stack(data = list(zinc = NA),
A = list(A.pred, 1),
effects = list(c(s.index, list(Intercept = 1)),
list(dist = meuse.grid$dist)),
tag = "meuse.pred")

#Join stack
join.stack <- inla.stack(meuse.stack, meuse.stack.pred)

```

Same in plain INLA (2)

Notes

```
#Fit model
form <- log(zinc) ~ -1 + Intercept + dist + f(spatial.field, model = spde)

m1 <- inla(form, data = inla.stack.data(join.stack, spde = meuse.spde),
           family = "gaussian",
           control.predictor = list(A = inla.stack.A(join.stack), compute = TRUE))
```

Note: We still have not compute predictions... and this is not too easy in plain INLA!!

When is `inlabru` easier to use

Notes

- spatial modeling.
 - point processes.
 - multiple likelihoods
 - when interested in spatial predictions
 - **inlabru** is also useful if one has non-linearities in the predictor η
 - born for ecological models (for example transect sampling) but used also in other fields

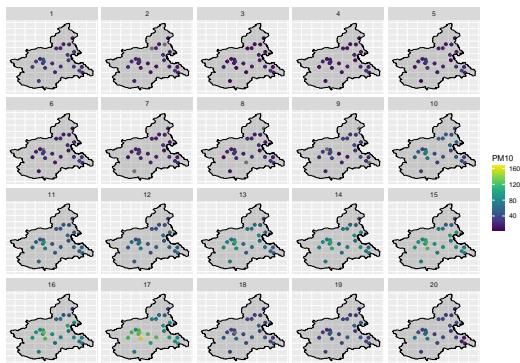
What cannot (at the moment) be done with inlabru

- Survival models

Space-time Modeling

Notes

Modeling PM10 concentration in Piemonte (Italy)



Notes

What have we learned Spatial models Gaussian Random Field models The SPDE approach Fitting spatial models Space-time Modeling Prediction in space

Modeling PM10 concentration in Piemonte (Italy)

Notes

Example from the Blangiardo & Cameletti book but simplified!

```

## Station.ID Date dem x y ws tmin HMX PREC EMI W10
## # 1 2005-10-01 95.2 469.45 4972.85 0.90 288.81 1294.6 0 10.05 28
## # 2 2005-10-01 164.1 423.48 4950.69 0.82 288.67 1139.8 0 18.74 22
## # 3 2005-10-01 242.9 490.71 4948.96 0.96 287.44 1404.0 0 6.28 17
## # 4 4 2005-10-01 149.9 437.36 4973.34 1.17 288.63 1042.4 0 29.35 25
## # 5 5 2005-10-01 405.0 426.44 5045.66 0.60 287.63 1038.7 0 32.19 20
# time logPM10
## # 1 1.332205
## # 2 1.3.091042
## # 3 1.2.833213
## # 4 1.3.218876
## # 5 1.2.995732

```

What have we learned	Spatial models	Gaussian Random Field models	The SPDE approach	Fitting spatial models	Space-time Modeling	Prediction in space
oooooooooooo	oooooo	oooooo	oooo	oooooooooooooooooooo	oooo•oooooooooooo	oooooooooooo

The model

We model the log PM10 concentration:

$$y_{it} \sim \mathcal{N}(\eta_{it}, \sigma_e^2)$$

$$\eta_{it} = \alpha + \beta_1 \text{dem}_i + \beta_2 \text{temp}_{it} + \omega_{it}$$

- y_{it} is the log-concentration at location i in time t
 - α is an intercept
 - β_1 and β_2 parameters of altitude and temperature
 - ω_{it} is the space-time residual

Space time residual model

A first order autoregressive process with spatially colored innovations

$$\omega_{it} = a\omega_{i(t-1)} + \xi_{it}$$

- $|a| < 1$ parameter of the AR1 process
 - ξ_{it} is a zero mean, temporally independent, Gaussian field with

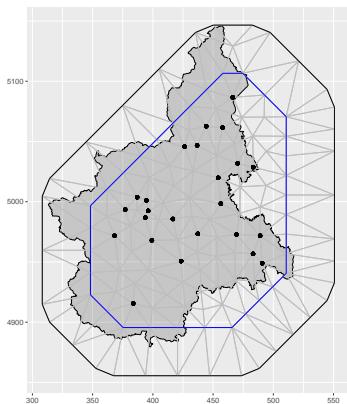
$$\text{Cov}(\xi_{it}, \xi_{ju}) = \begin{cases} 0, & \text{if } t \neq u \\ \mathcal{C}(h), & \text{if } t = u \end{cases}$$

- h distance between locations i and j - $\mathcal{C}(h)$ is a Matern correlation function

Notes

What have we learned Spatial models Gaussian Random Field models The SPDE approach Fitting spatial models Space-time Modeling Space-time Models Prediction in space

Building the model: mesh



Notes

What have we learned Spatial models Gaussian Random Field models The SPDE approach Fitting spatial models Space-time Modeling Prediction in space

Building the model: Prior for the spde model

Notes

```
spde = inla.spde2.pcmatern(mesh = mesh,
                           prior.range=c(100, 0.5),
                           prior.sigma=c(1, 0.1))
```

- $\text{Prob}(\rho < 100 \text{ Km}) = 0.5$
 - $\text{Prob}(\sigma > 1) = 0.1$

Implement the model

```
# make the data a spatial object
coordinates(df) = c("x","y")

# model component
cmp = ~ Intercept(1) +
  SPDE(coordinates, model = spde,
        group = time, control.group = list(model = "ar1")) +
  dem(dem, model = "linear") +
  temp(temp, model = "linear")

# likelihood
lik = like(formula = logPM10 ~ Intercept + SPDE + dem + temp ,
           family = "gaussian",
           data = df)

# fit the model
fit = bru(cmp, lik,
          options = list(verbose = F,
                         bru_max_iter = 1,
                         inla.mode = "experimental"))
```

Notes

Model Parameters

Notes

Fixed effects:

```

##               mean        sd   0.025quant  0.975quant
## Intercept -1.438103e+01 9.2183351517 -32.139029373 4.0476156436
## dem        -7.283307e-04 0.0008432533 -0.002398673 0.0009354501
## temp       6.250900e-02 0.0320760645 -0.001679595 0.1243716815

```

Hyperparameters of the random field:

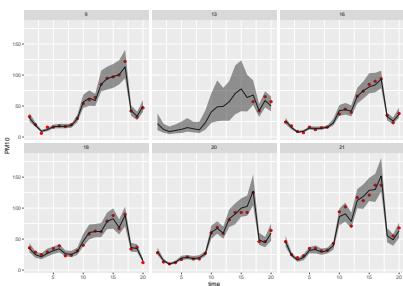
```

##                                     mean      sd 0.025quant
## Precision for the Gaussian observations 41.8579579 7.23518899 29.185249
## Range for SPDE 129.1799687 16.96397904 101.0510932
## Stddev for SPDE 0.7511733 0.09663733 0.5913784
## GroupRho for SPDE 0.0105269 0.02594638 0.8555876
##                                         0.9795555
## Precision for the Gaussian observations 57.6625010
## Range for SPDE 167.2757328
## Stddev for SPDE 0.9796968
## GroupRho for SPDE 0.9557794

```

Notes

Prediction at station locations



Notes

Prediction is space

Notes

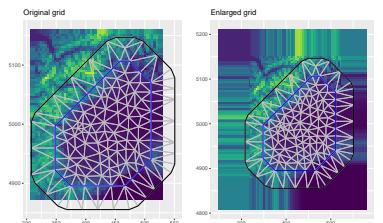
- To do this we need to have the covariates over the space of interest.
 - They should cover the whole mesh

Expanding the covariates: the `bru_fill_missing`

Notes

Expanding the covariates: the bru_fill_missing

Notes



Note: You can do this with your favourite method!

Notes

Prediction in space

What have we learned Spatial models Gaussian Random Field models The SPDE approach Fitting spatial models Space-time Modeling Prediction in space

Prediction in space

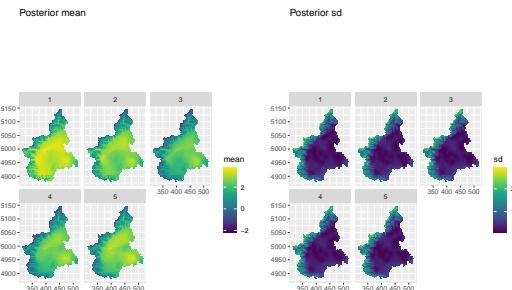
Notes

```
# Create a space time grid
pxi = pixels(mesh, nx = 200, ny = 200, mask = shape)
ips2 <- ipoints(domain = c(1:5), name = "time")
pxl_time <- cbind(cprod(pxi, ips2), data.frame(temp = 0))

pred_space = predict(fit, pxi_time,
- Intercept + SPDE +
  dem_eval(inlabru:::eval_SpatialDF(dem_large,
  .data...)) +
  temp_eval(inlabru:::eval_SpatialDF(temp_large,
  .data...,
  selector = "time")))
```

Prediction in space

Notes



Samples form the fitted model

Notes

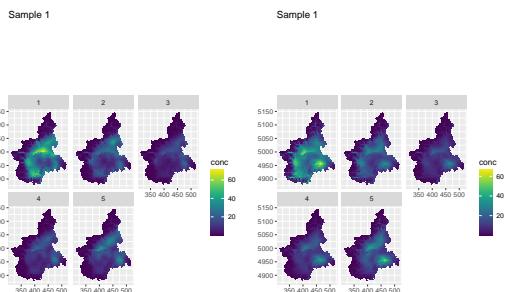
It is also possible to generate from the fitted model

```
sim_space = generate(fit, pxl_time,
- exp(Intercept + SPDE +
  dem_eval(inlabru:::eval_SpatialDF(dem_large,
  .data.)) +
  temp_eval(inlabru:::eval_SpatialDF(temp_large,
  .data.,
  selector = "time"))),
n.samples = 2)
```

This can be useful as the posterior means are always smoother than the “real” field

Samples form the fitted model

Notes



Some concluding remark

Notes

- The functions `predict()` and `generate()` in `inlabru` use the INLA function `inla.posterior.sample()` internally.
- `inlabru` is a wrapper around INLA so all the internal computations are identical!
- `inlabru` makes handling of spatial object a lot easier
- Transition from `sp` to `sf/terra` some changing can be expected

Notes

Thank you for your attention!

If you have any doubts or questions, please write :
sara.martino@math.ntnu.no



Notes

Notes
